

Spark Practice

1. Overview
2. Spark Applications
3. Resilient Distributed Datasets
4. Datasets
5. Architecture
6. Operations
7. Tools
8. Conclusion

1. Overview

- Apache software, from UC Berkeley
 - Compatible with Hadoop (HDFS)
- Extension of MapReduce for two classes of analytics applications
 - Iterative processing (machine learning, graphs)
 - Interactive data mining (R, Excel, Python)
- Major performance improvement (up to 100*)
 - In-memory processing
 - Optimization of the task graph
- Major usability improvement
 - APIs for Java, Python, R and Scala (functional extension of Java)
 - Interactive use from a Scala interpreter
- Major adoption from industry
 - Databricks: a successful startup from UC Berkeley
 - Will replace MapReduce

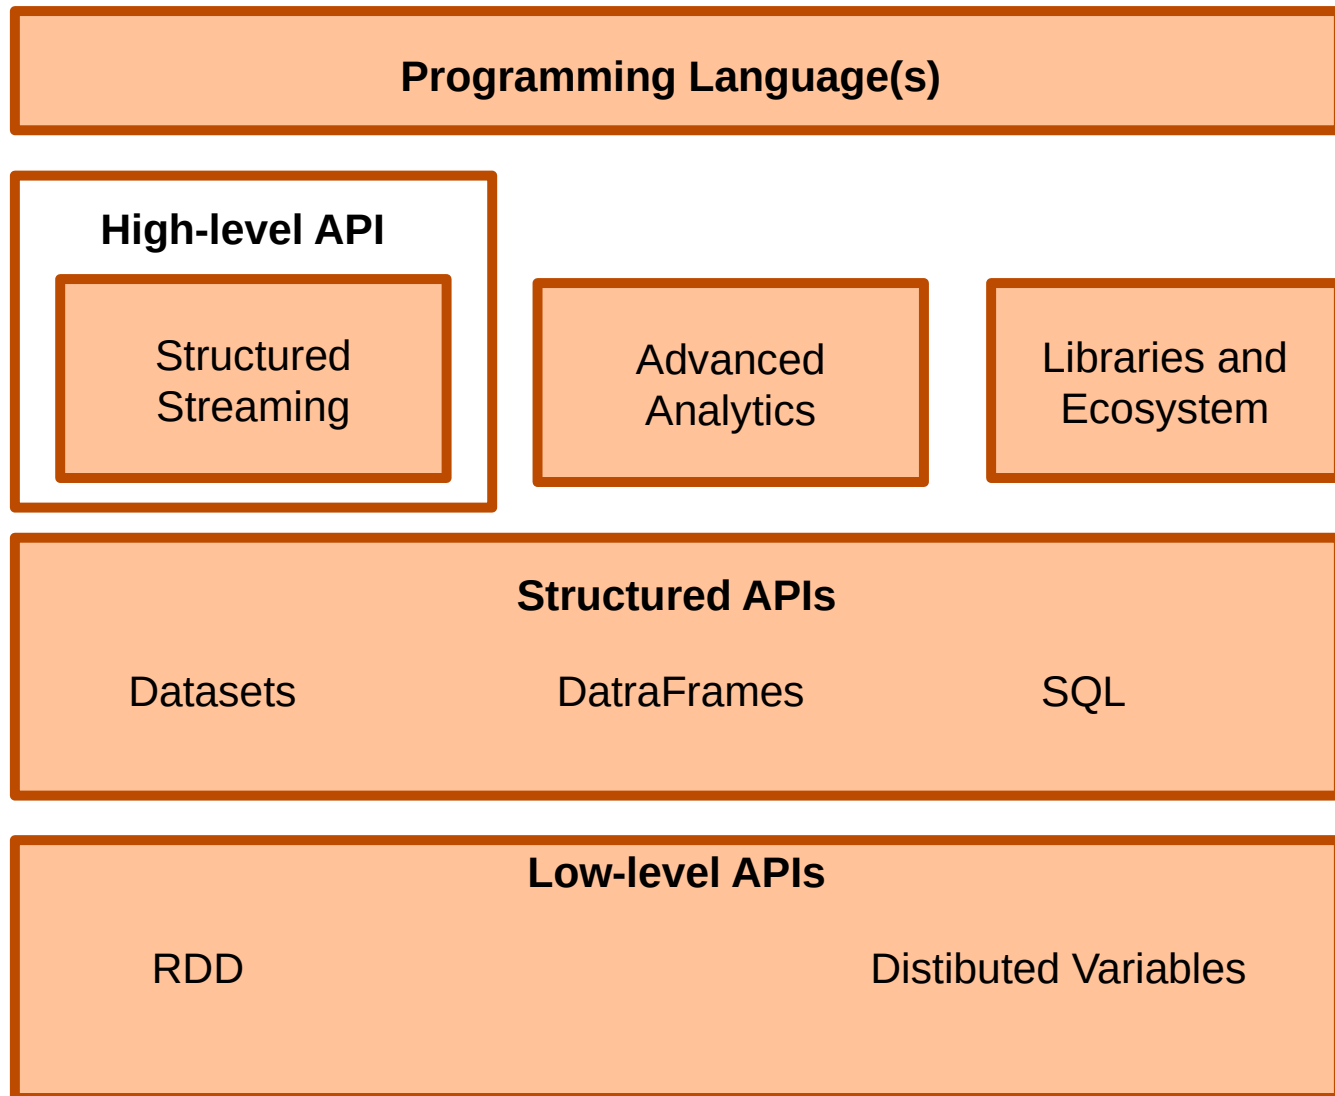
What is Spark?

- Unified computing engine and set of libraries for parallel data processing on computer cluster
 - Unified platform for writing big data applications
 - Computing engine:
 - Spark's focus on computation
 - Loading data from storage systems and performing computations on it
 - Using Spark with a wide variety of persistent storage systems
 - Libraries: unified API for common data analysis tasks
 - SQL and structured data (Spark SQL)
 - Machine learning (Mlib)
 - Stream processing (Spark Streaming and the newer Structured Streaming)
 - Graph analytics (GraphX)

Spark Model

- **Concept : Resilient Distributed Dataset (RDD)**
 - A collection of objects distributed in a cluster
 - Built from parallel transformations (map, filter, etc.)
 - Can be maintained in memory for efficient processing
 - Preserves the nice properties of MapReduce
 - Fault-tolerance, data locality, scalability
- **An RDD comes with its provenance information**
 - How it has been produced from other RDDs
 - Useful to rebuild an RDD that has been lost after a crash
- **The user can control**
 - Data persistence (disk or RAM)
 - Data partitioning (hashing, range, [k , v])
- **Operators: transformations and actions**

Spark Toolset (1)



Spark Toolset (2)

- Widely used programming languages (Python, Java, Scala, and R)
- Structured streaming: high-level API for stream processing
 - Allows to rapidly and quickly extract value out of streaming systems with virtually no code change
 - Easy to conceptualize
- Advanced analytics
 - Techniques aimed at solving the core problem of deriving insights
 - Making predictions or recommendations based on data
- Ecosystem
 - Packages and tools created by the community
 - [Spark-packages.org](http://spark-packages.org)

Spark Toolset (3)

- Libraries
 - Unified API for data analysis
 - Provides more and more types of functionality
 - Hundreds of open source external libraries
- Structured APIs for manipulating all sorts of data such as
 - Unstructured log files
 - Semi-structured CSV files
 - Highly structured Parquet files
- Low-level APIs
 - Manipulating distributed data (RDDs)
 - Distributing and manipulating distributed shared variables
 - Broadcast variables
 - Accumulators

2. Spark Application (1)

- Set of Spark jobs defined by one Spark context in the driver program
- Two components
 - Driver process: runs the main() function, sits on a node in the cluster
 - Maintains information about the Spark application
 - Responds to a user's program or input
 - Analyzes, distributes, and schedules work across the executors
 - Executor: carrying out the work that the driver assigns them
 - Executes code assigned by the driver
 - Reports the state of computation on that executor back to the driver node

Spark Application (2)

- SparkSession is a driver process that enables to control Spark Application

In Scala:

```
res0: org.apache.spark.sql.Session =  
org.apache.spark.sql.Session@...
```

In Python:

```
<pyspark.sql.session.Session at 0x7efda4c1ccd0>
```

- The SparkSession instance is the way Spark executes user-defined manipulations across the cluster

Spark Application (3)

- SparkContext object is the driver program that communicates with the appropriate cluster manager to run the tasks.

```
# in Python
```

```
from pyspark.sql import Row
schema = df.schema
newRows = [
    Row("New Country", "Other Country", 5L),
    Row("New Country 2", "Other Country 3", 1L)
]
parallelizedRows = spark.sparkContext.parallelize(newRows)
newDF = spark.createDataFrame(parallelizedRows, schema)
```

```
# in Python
```

```
df.union(newDF)\
.where("count = 1")\
.where(col("ORIGIN_COUNTRY_NAME") != "United States")\
.show()
```

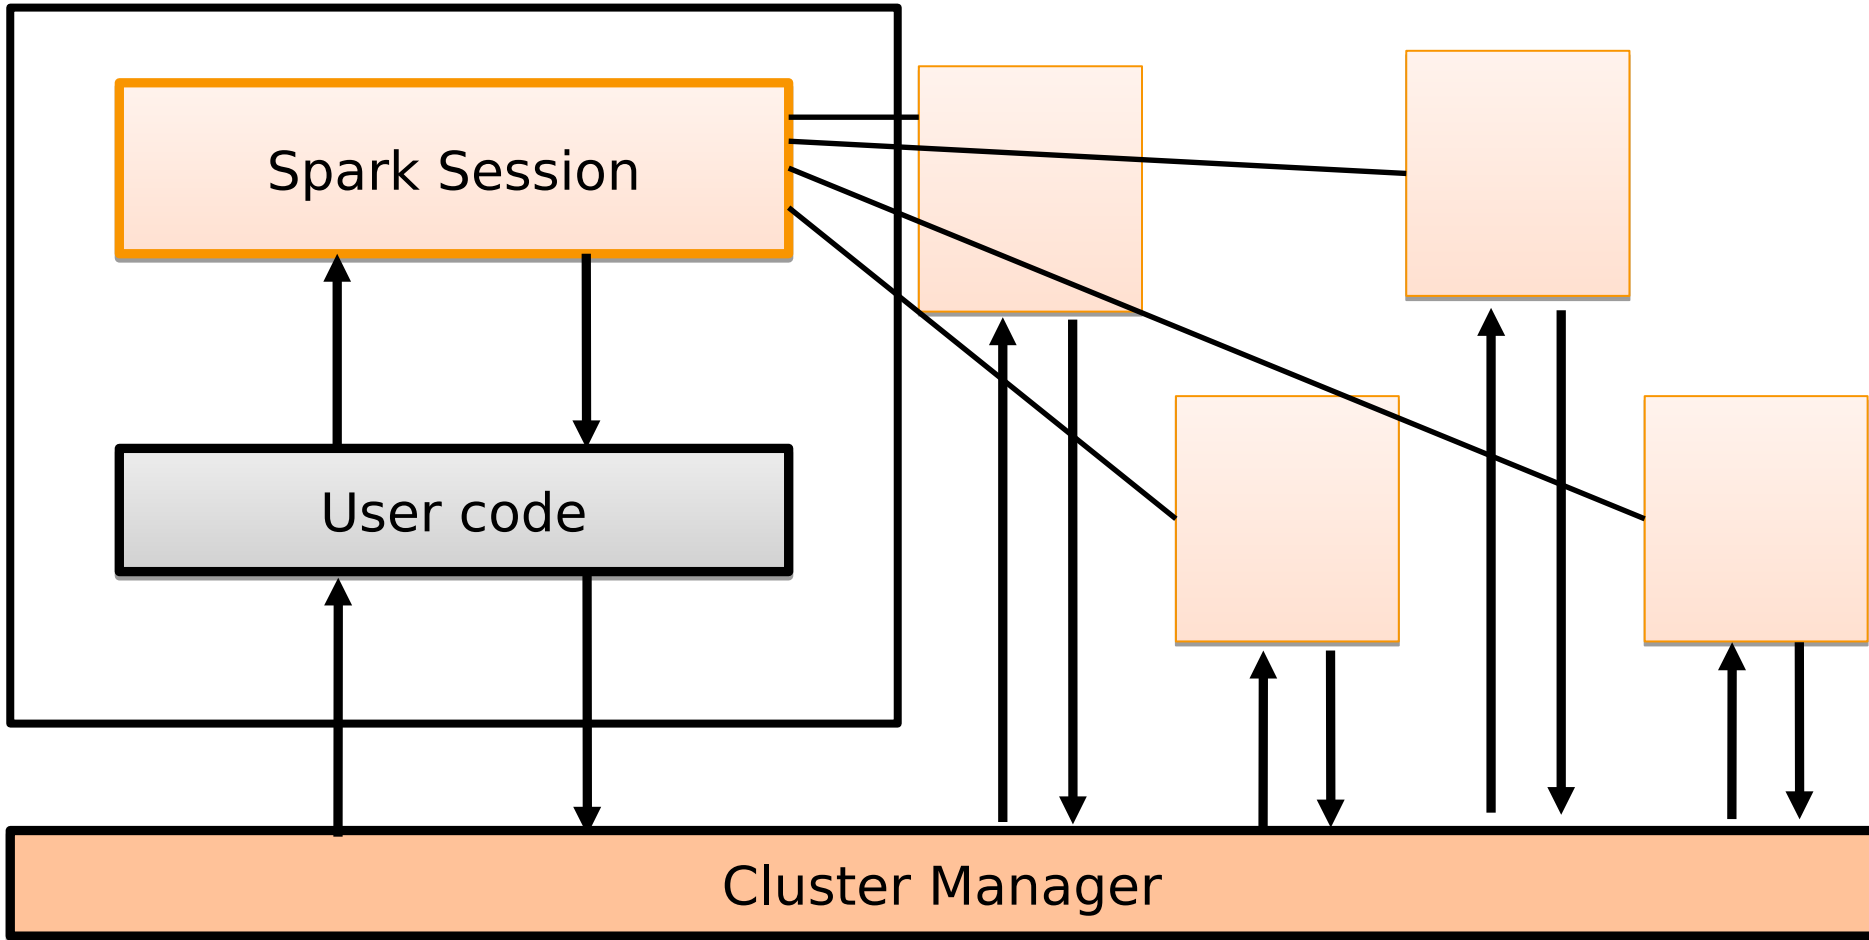
Giving the output of:

```
+-----+-----+-----+
|DEST_COUNTRY_NAME|ORIGIN_COUNTRY_NAME|count|
+-----+-----+-----+
| United States   | Croatia            | 1    |
```

Spark Application Architecture

Driver Process

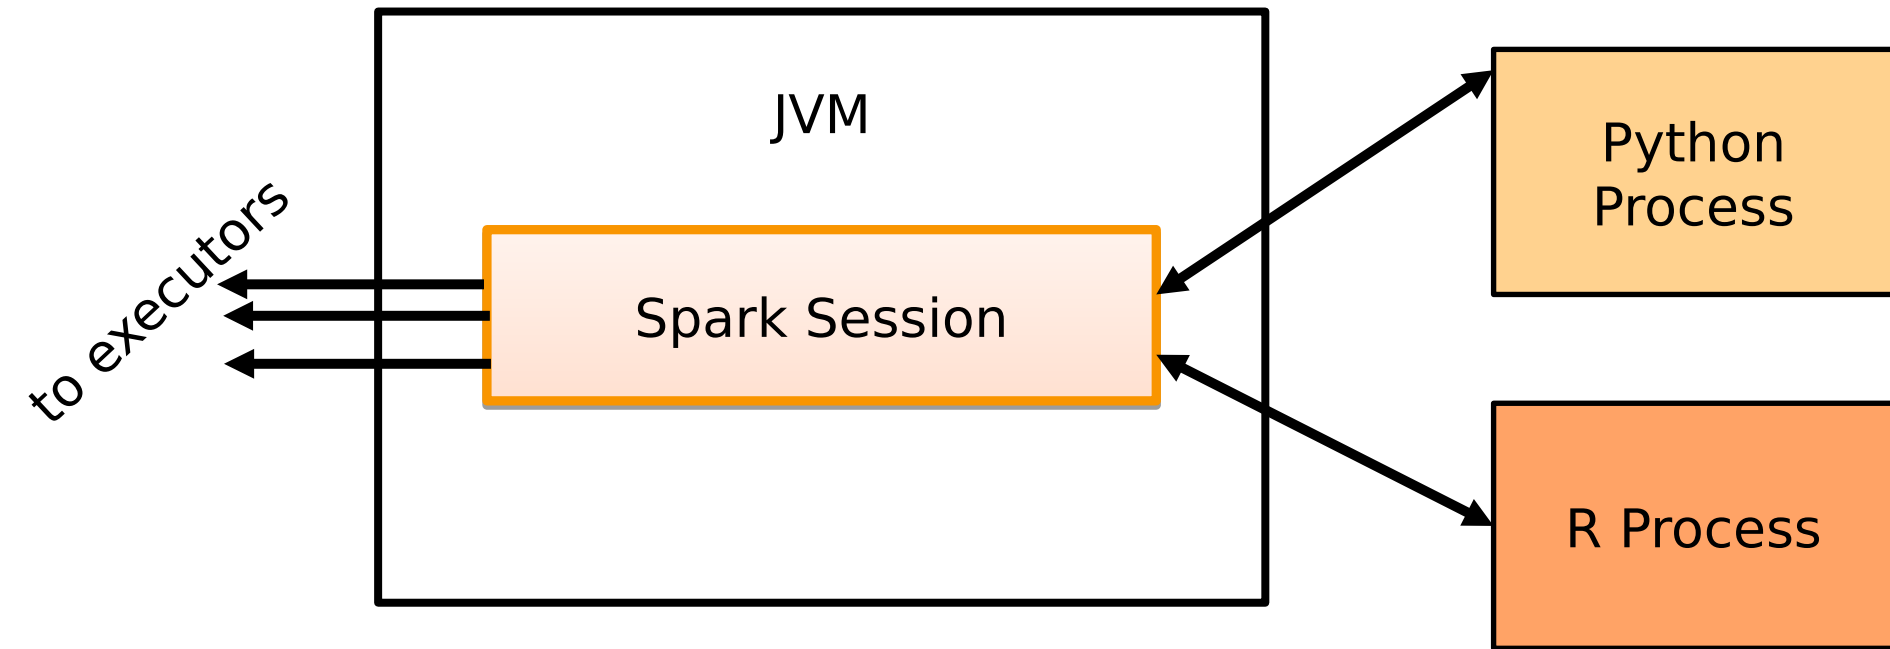
Executors



Spark Language

- Spark code uses several programming languages
 - Scala (Spark's default language)
 - Java
 - Python (supports nearly all constructs that Scala supports)
 - SQL (supports a subset of the ANSI SQL 2003 standard)
 - R Spark uses Spark core (SparkR) and R community-driven package (sparklyr)
- Spark has some core concepts in every language
 - Translating concepts into Spark code
 - Runs Spark code on a cluster of machines

Relationship between SparkSession and Spark's Language API



3. Resilient Distributed Datasets

- Resilient Distributed Dataset (RDD)
 - In memory (possibly cached)
 - Read-only collection of objects
 - Partitioned across machines of a cluster
 - Lineage: allows partitions to be rebuilt when failure

RDD API

- Set of partitions (splits)
 - In memory (possibly cached)
- List of dependencies
- Function to compute a partition
- Preferred locations
- Partitioner

RDD Construction

- From a file (e.g., in HDFS)
`val rdd = sc.textFile("hdfs://...")`
- By parallelizing a collection
`val data = Range(0, 100)`
`val rdd = sc.parallelize(data)`
- By transforming another RDD
`val rdd2 = rdd1.filter(...)`

Operations

- Transformations

- $\text{RDD}(s) \Rightarrow \text{RDD}$
- Lazily applied
 - Only computed if result is returned
- Chain of transformations recomputed each time
 - Except if persisted with `cache()`

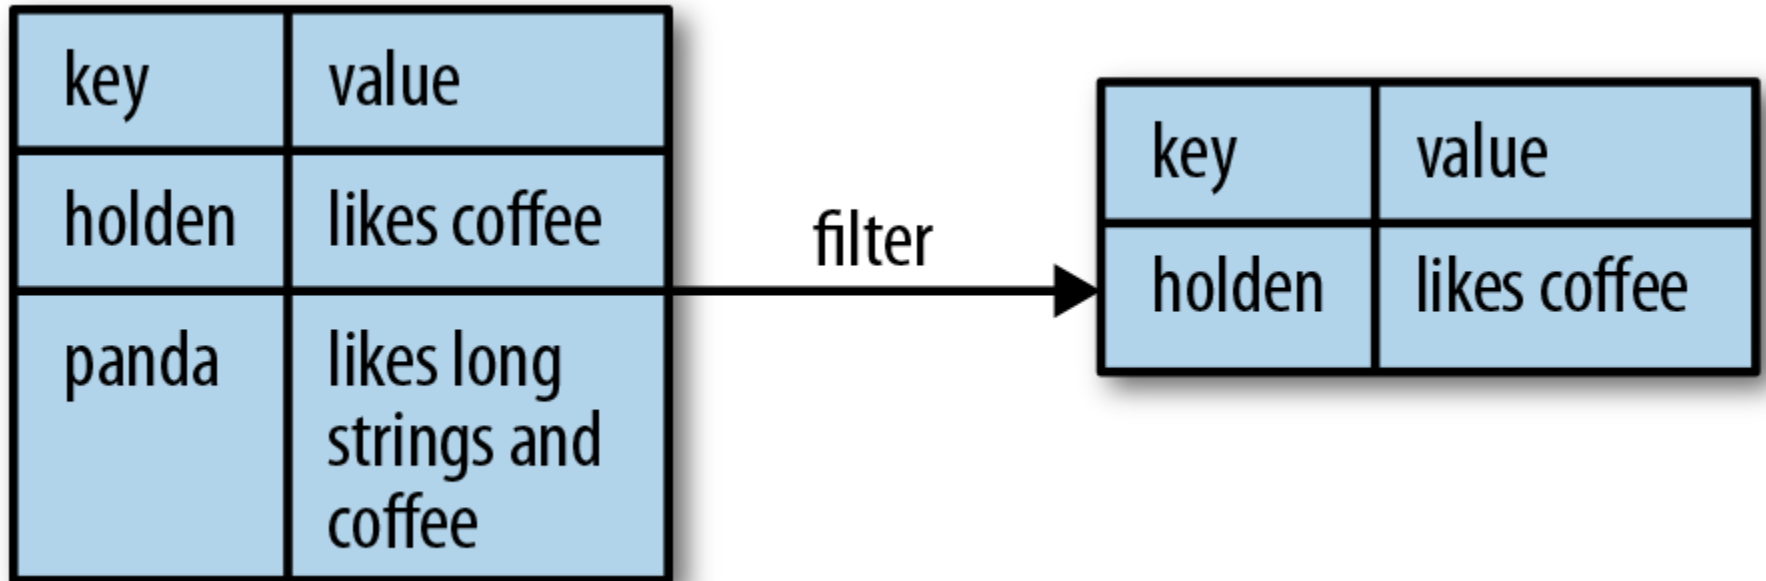
- Actions

- $\text{RDD} \Rightarrow \text{value}$

Transformations: simple, no shuffle

Transformation	Types
map(f: T => U)	RDD[T] \Rightarrow RDD[U]
filter(f: T => boolean)	RDD[T] \Rightarrow RDD[T]
flatMap(f: T => Seq[U])	RDD[T] \Rightarrow RDD[U]
mapPartitions(f: Seq[T] => Seq[U])	RDD[T] \Rightarrow RDD[U]
mapPartitionsWithIndex(f: (int, Seq[T]) => Seq[U])	RDD[T] \Rightarrow RDD[U]
sample()	RDD[T] \Rightarrow RDD[T]
sortBy(f: T => K)	RDD[T] \Rightarrow RDD[T]
keyBy(f: T => K)	RDD[T] \Rightarrow RDD[(K,T)]

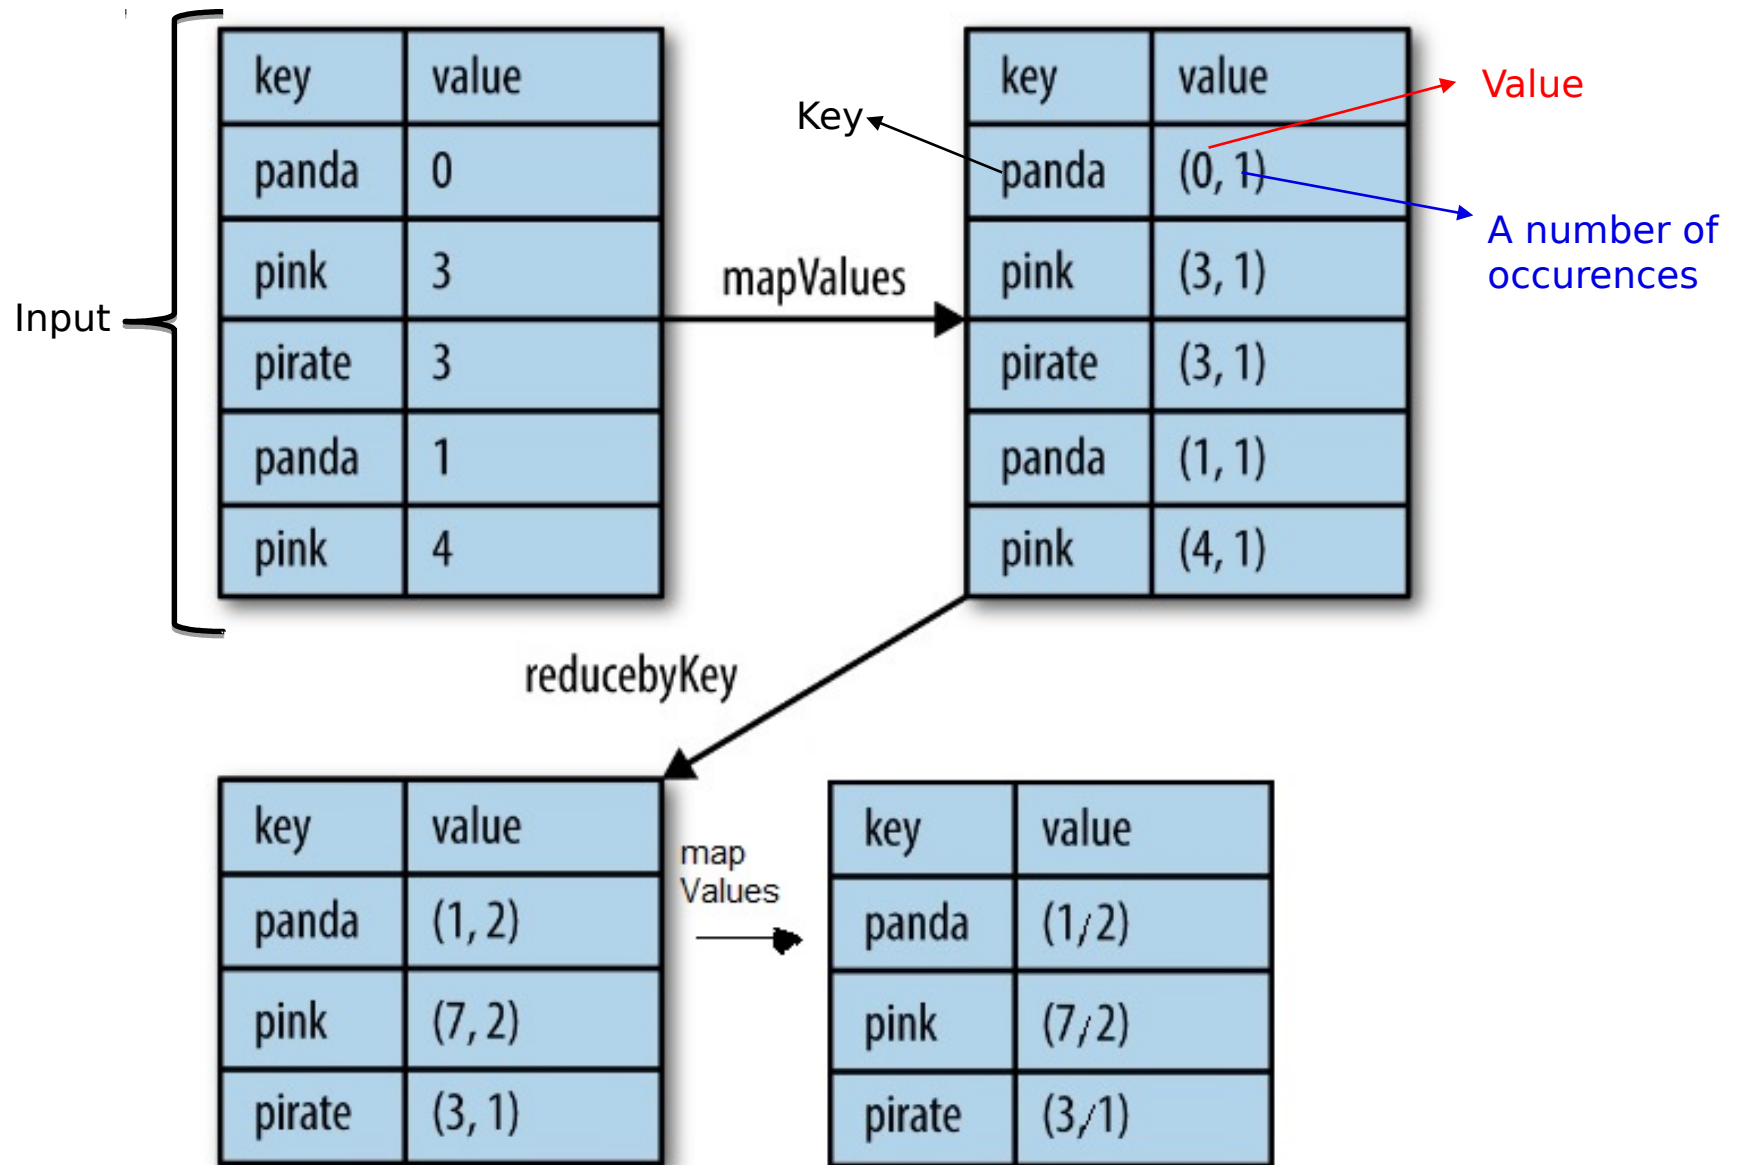
Transformations: simple, no shuffle



transformations: key-value, no shuffle

Transformation	Types
mapValues(f: V => U)	RDD[(K,V)] \Rightarrow RDD[(K,U)]
flatMapValues(f: V => Seq[U])	RDD[(K,V)] \Rightarrow RDD[(K,U)]
sampleByKey()	RDD[(K,V)] \Rightarrow RDD[(K,V)]

Transformations: key-value, shuffle



Transformations: key-value, shuffle

Transformation	Types
groupByKey()	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[(K, \text{Seq}[V])]$
reduceByKey(f: (V,V) => V)	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[(K,V)]$
aggregateByKey(zero: U) (f: (U,V) => U, g: (U,U) -> U)	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[(K,V)]$
combineByKey(f: V => C, g: (C, V) => C, h: (C, C) => C)	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[(K,C)]$
foldByKey(zero: V)(f: (V, V) => V)	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[(K,V)]$
groupByKey()	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[(K, \text{Seq}[V])]$
partitionBy(part: K => int)	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[(K,V)]$
reduceByKey(f: (V, V) => V)	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[(K,V)]$
keys()	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[K]$
values()	$\text{RDD}[(K,V)] \sqsubseteq \text{RDD}[V]$

Transformations: join & grouping

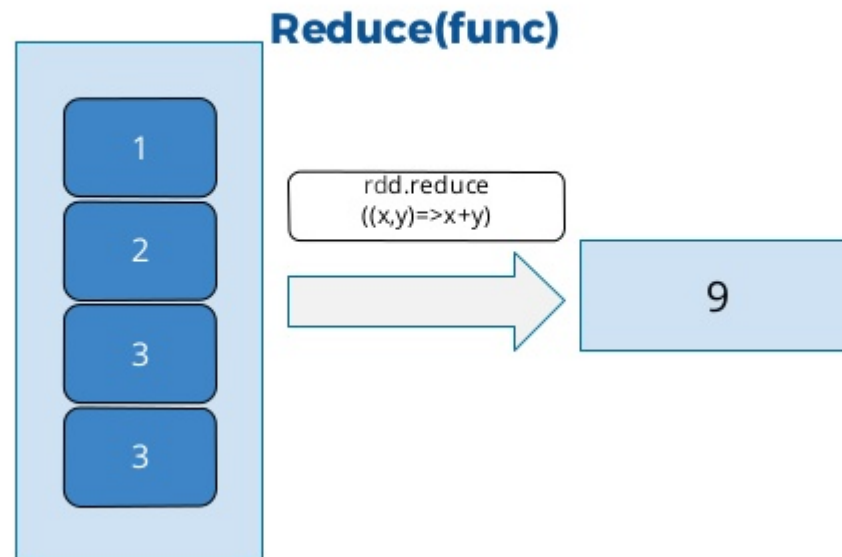
Transformation	Types
join(RDD[(K, W)])	(RDD[(K,V)], RDD[(K,W)]) \sqsubseteq RDD[(K,(V,W))]
leftOuterJoin(RDD[(K, W)])	(RDD[(K,V)], RDD[(K,W)]) \sqsubseteq RDD[(K,(V,W))]
rightOuterJoin(RDD[(K, W)])	(RDD[(K,V)], RDD[(K,W)]) \sqsubseteq RDD[(K,(V,W))]
fullOuterJoin(RDD[(K, W)])	(RDD[(K,V)], RDD[(K,W)]) \sqsubseteq RDD[(K,(V,W))]
cogroup(RDD[(K, W)])	(RDD[(K,V)], RDD[(K,W)]) \sqsubseteq RDD[(K,(Seq[V],Seq[W]))]
subtractByKey(RDD[(K, W)])	(RDD[(K,V)], RDD[(K,W)]) \sqsubseteq RDD[(K,V)]

Actions

Action	Types
count()	RDD[T] \sqsubseteq Long
reduce(f: (T, T) => T)	RDD[T] \sqsubseteq T
treeReduce(f: (T, T) => T)	RDD[T] \sqsubseteq T
aggregate(zero: U) (f: (U,T) => U, g: (U,U) => U)	RDD[T] \sqsubseteq U
treeAggregate(zero: U) (f: (U,T) => U, g: (U,U) => U)	
first()	RDD[T] \sqsubseteq T
take()	RDD[T] \sqsubseteq Array[T]
takeSample()	RDD[T] \sqsubseteq Array[T]
max()	RDD[T] \sqsubseteq T

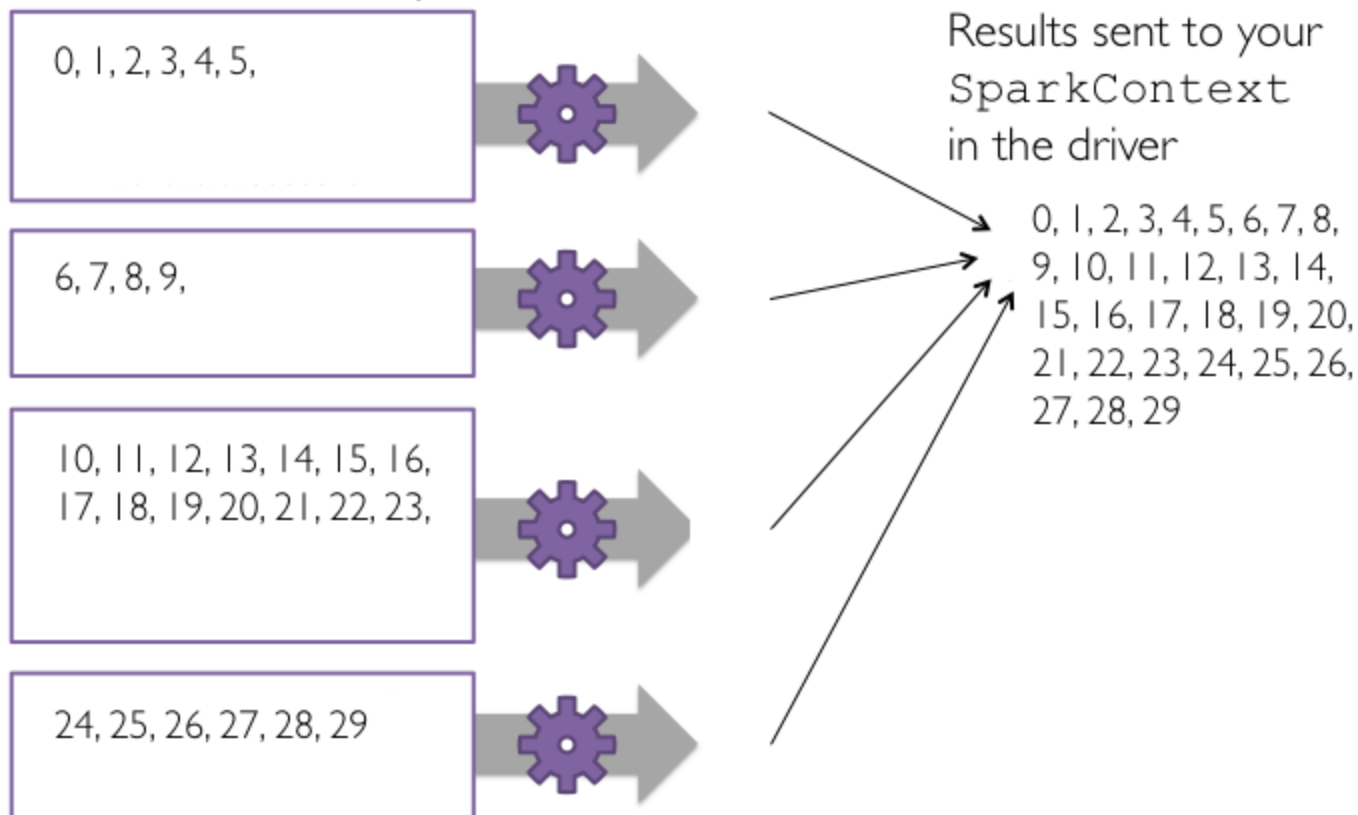
Actions

Actions



Actions

`collect()` : Gathers the entries from all partitions into the driver



Actions

Action	Types
countByKey()	RDD[(K,V)] \Rightarrow Map[K, Long]
collectAsMap()	RDD[(K,V)] \Rightarrow Map[K, V]
lookup(K)	RDD[(K,V)] \Rightarrow Seq[K, V]

Example : wordcount

```
val textFile = spark.textFile("hdfs://input/path")  
  
val words = textFile.flatMap(line => line.split("\\s+"))  
val oneCounts = words.map(word => (word, 1))  
val counts = oneCounts.reduceByKey(_ + _)  
  
val counts.saveAsTextFile("hdfs://output/path")
```

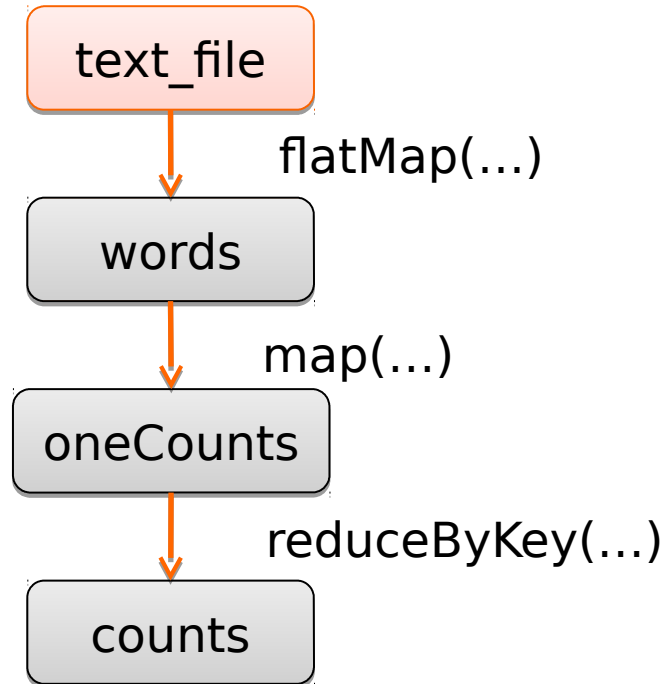
Example : wordcount

```
val textFile = spark.textFile("hdfs://input/path")

val counts = textFile.flatMap(line => line.split("\\s+"))
                        .map(word => (word, 1))
                        .reduceByKey(_ + _)

val counts.saveAsTextFile("hdfs://output/path")
```

Example : wordcount



4. Datasets

- Foundational type of the structured APIs
- A strictly Java Virtual Machine language feature that works only with Scala and Java
- When to use Datasets?
 - When the operation(s) you would like to perform cannot be expressed using dataframe manipulation
 - When you want or need type-safety
 - The cost of performance

Creating Datasets

In Java: Encoders

```
import org.apache.spark.sql.Encoders;
public class Flight implements Serializable{
    String DEST_COUNTRY_NAME;
    String ORIGIN_COUNTRY_NAME;
    Long DEST_COUNTRY_NAME;
}
Dataset<Flight> flights = spark.read
    .parquet("/data/flight-data/parquet/2010-
summary.parquet/")
    .as(Encoders.bean(Flight.class));
```


Creating Datasets

In Scala: Case Classes

```
case class Flight(DEST_COUNTRY_NAME: String,  
ORIGIN_COUNTRY_NAME: String, count: BigInt)  
val flightsDF = spark.read  
  .parquet("/data/flight-data/parquet/2010-  
summary.parquet/")  
val flights = flightsDF.as[Flight]
```

Creating Datasets

- To create Datasets in Scala, you define a Scala case class. A case class is a regular class
- That has the following characteristics:
 - Immutable
 - Decomposable through pattern matching
 - Allows for comparison based on structure instead of reference
 - Easy to use and manipulate

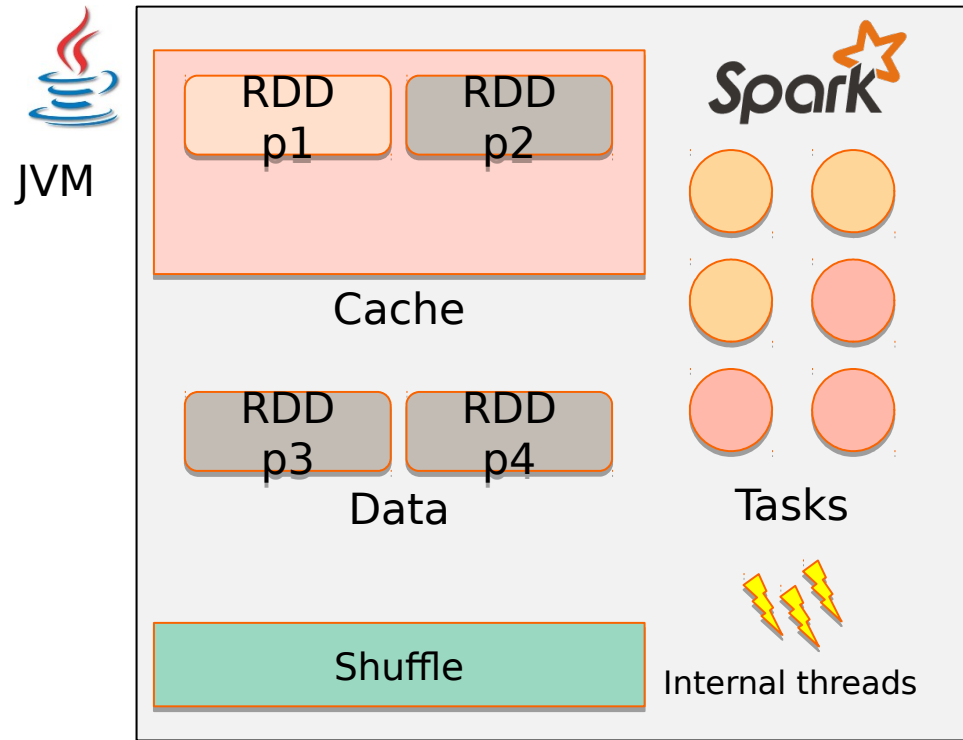
Operations

- Actions
 - Collect
 - Take
 - Count
- Transformations
 - Filtering
 - Mapping
 - Mapping example

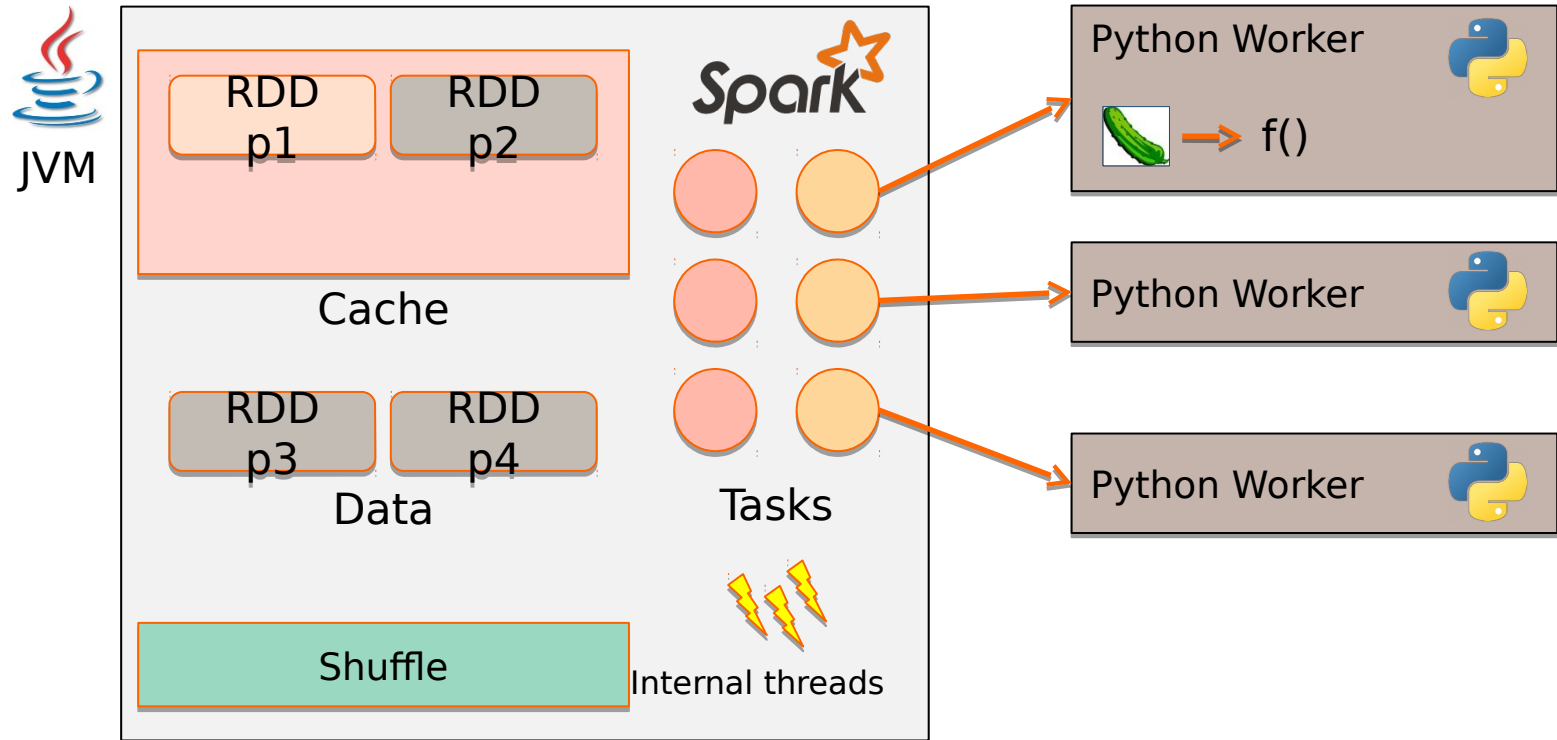
```
val destinations = flights.map(f =>  
  f.DEST_COUNTRY_NAME)
```

- Joins
- Grouping and Aggregations

5. Executor Architecture



Executor with PySpark



Deployment

- Roles

- Coordinator \Rightarrow Driver (SparkContext)
- Executor

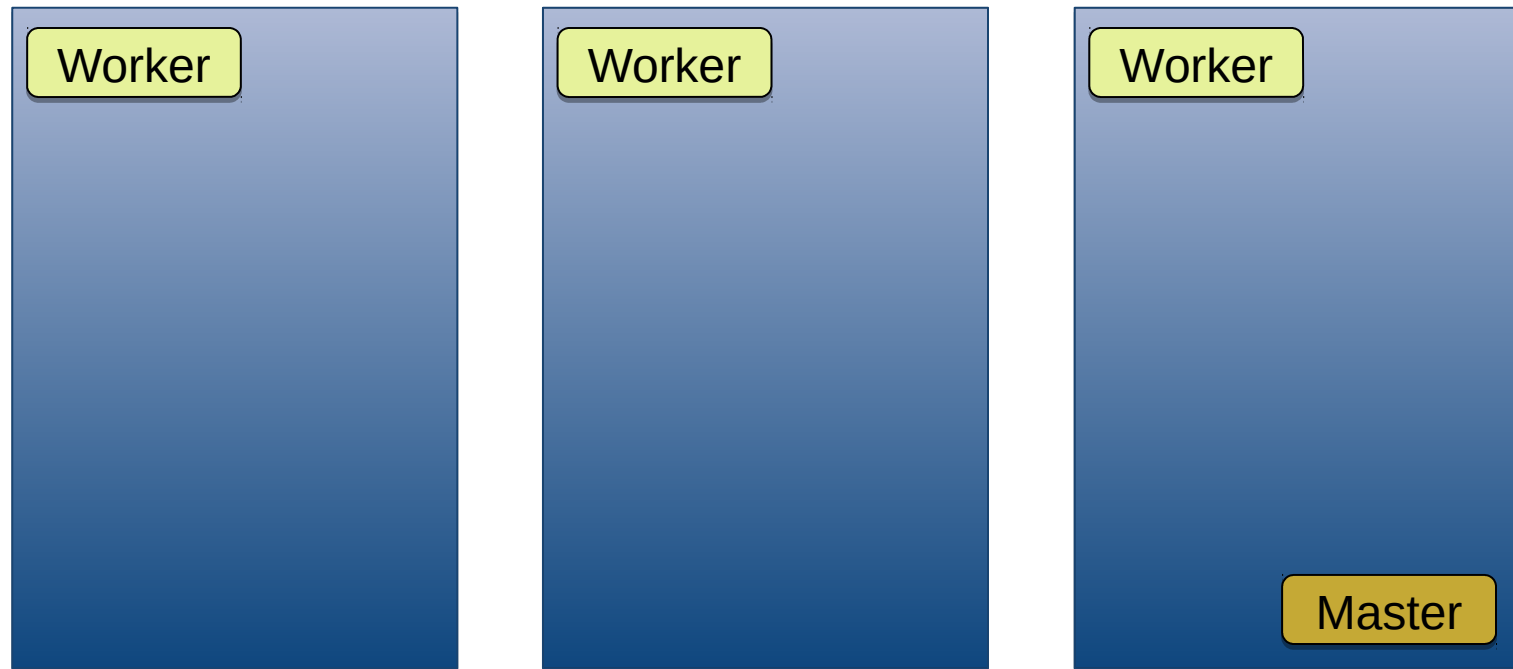
- Deployment

- Local
 - One executor (driver inside)
- Distributed
 - Driver + several executors

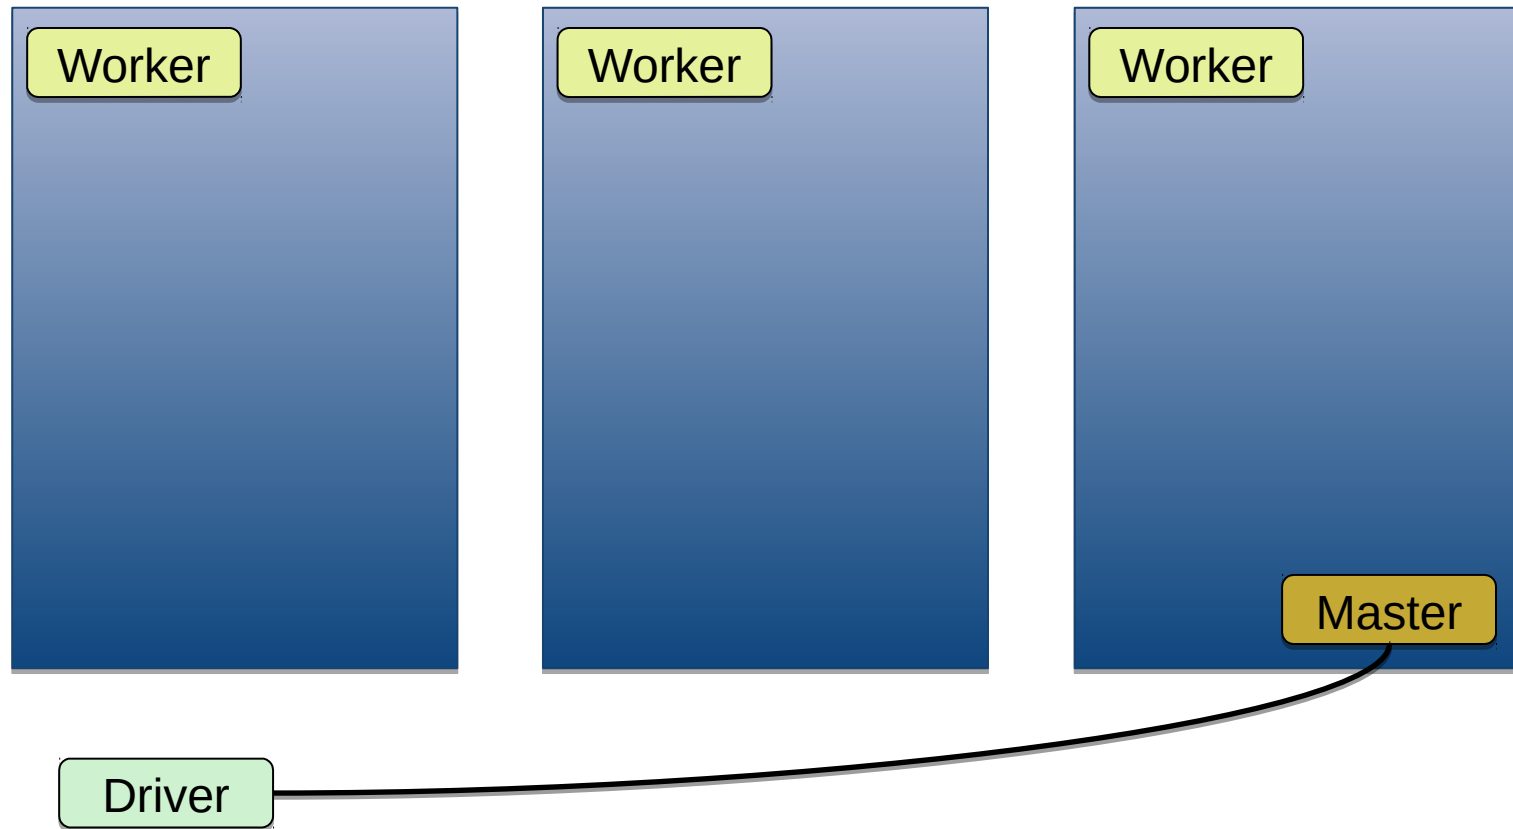
- Cluster manager maintains a cluster of machines that runs Spark applications

- Standalone
- YARN
- Mesos

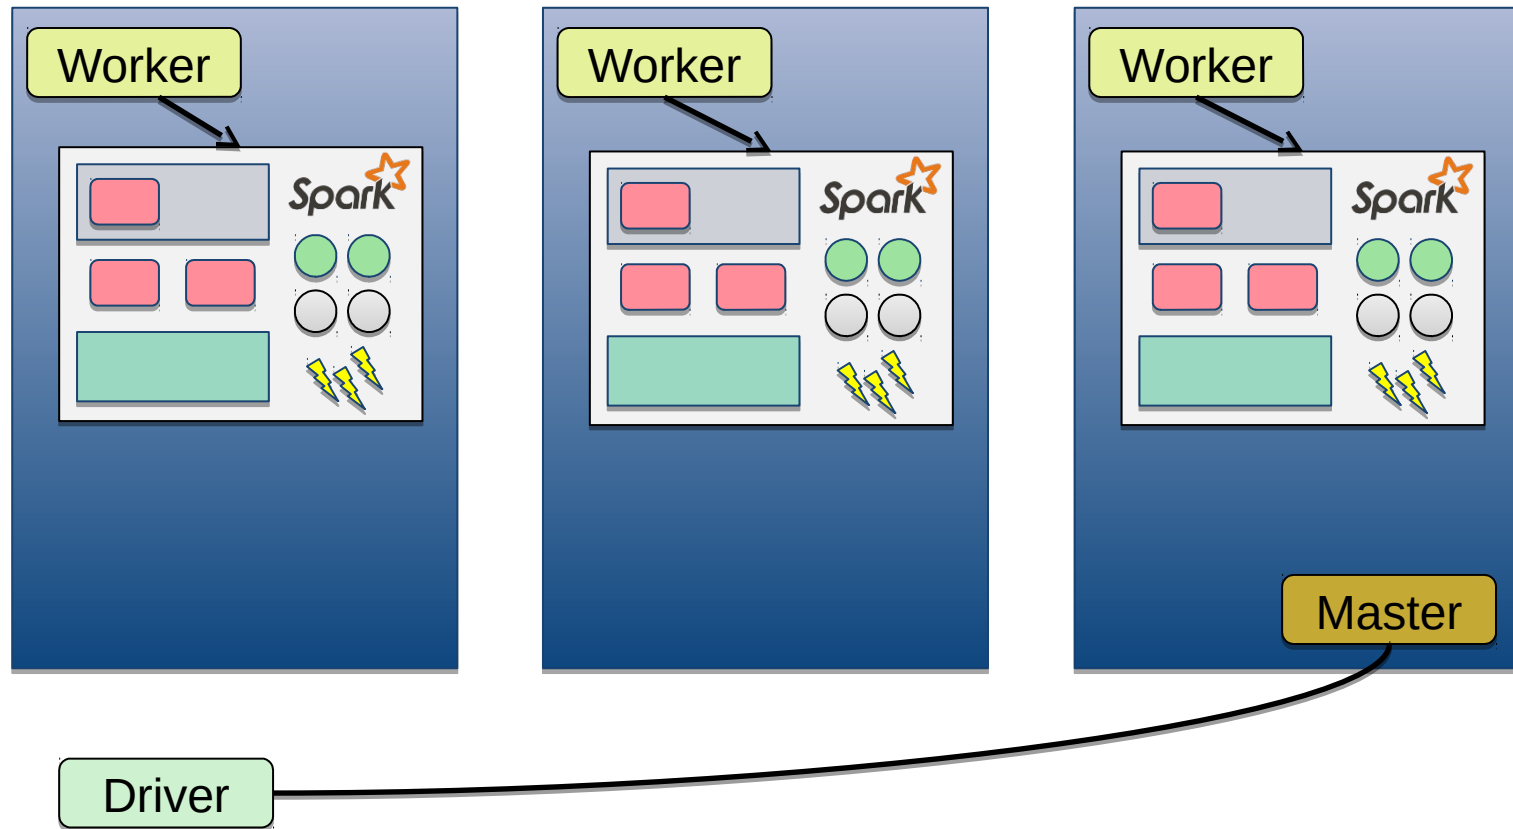
Spark Standalone



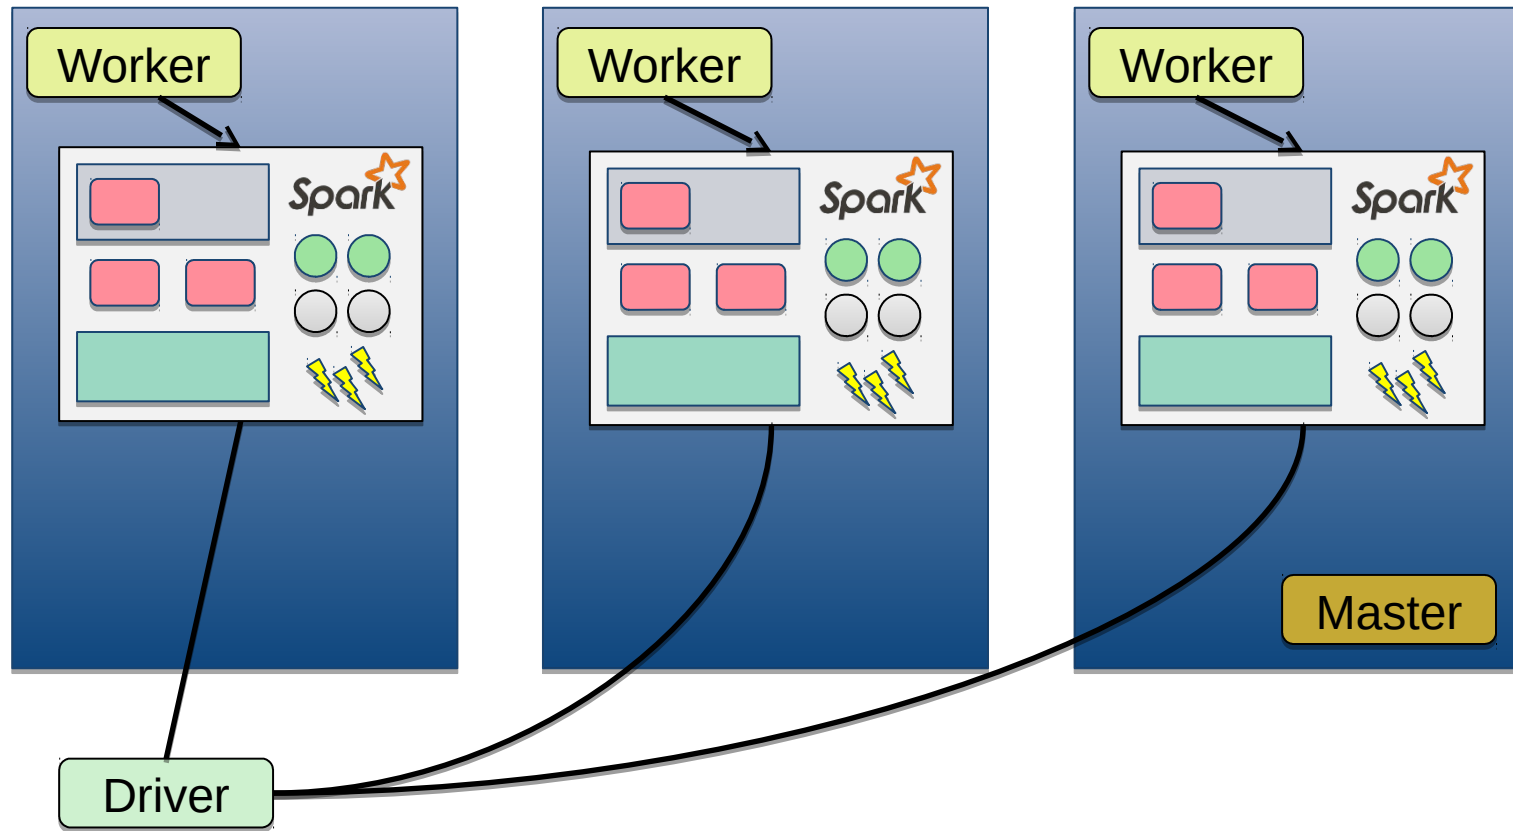
Spark Standalone



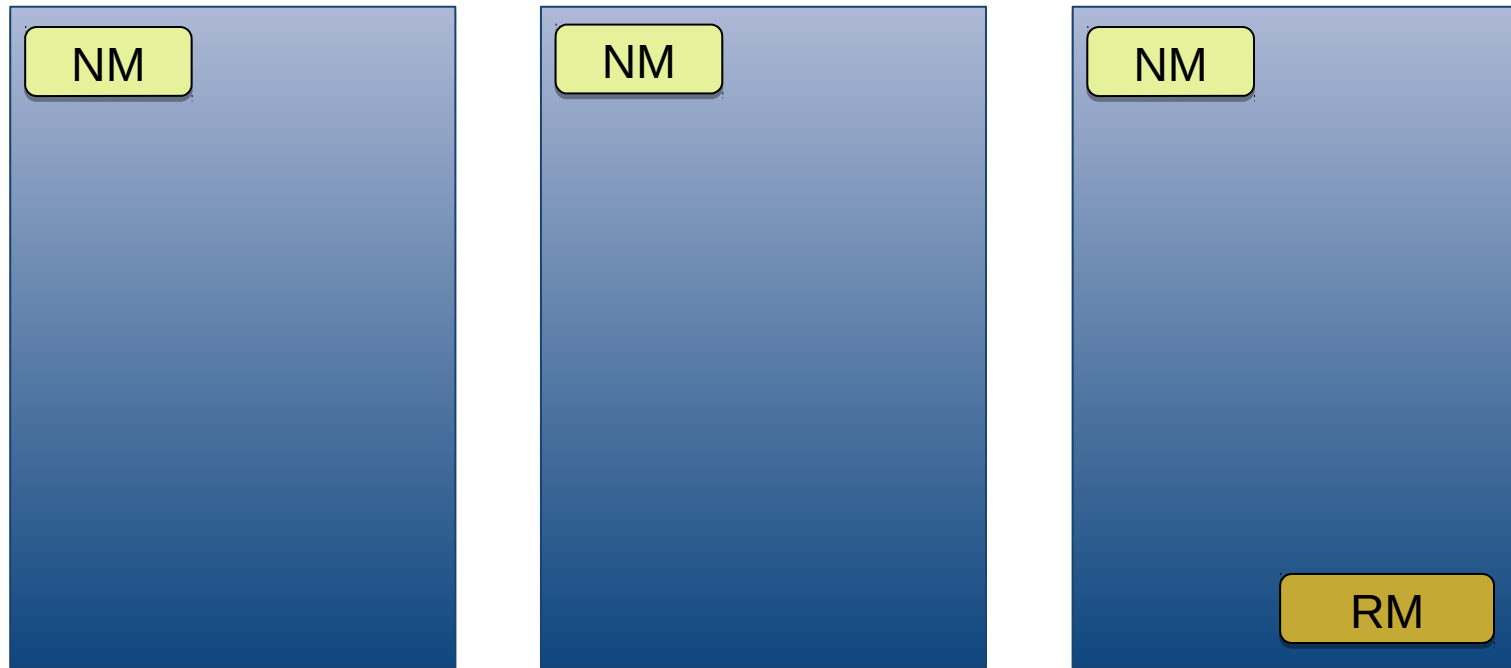
Spark Standalone



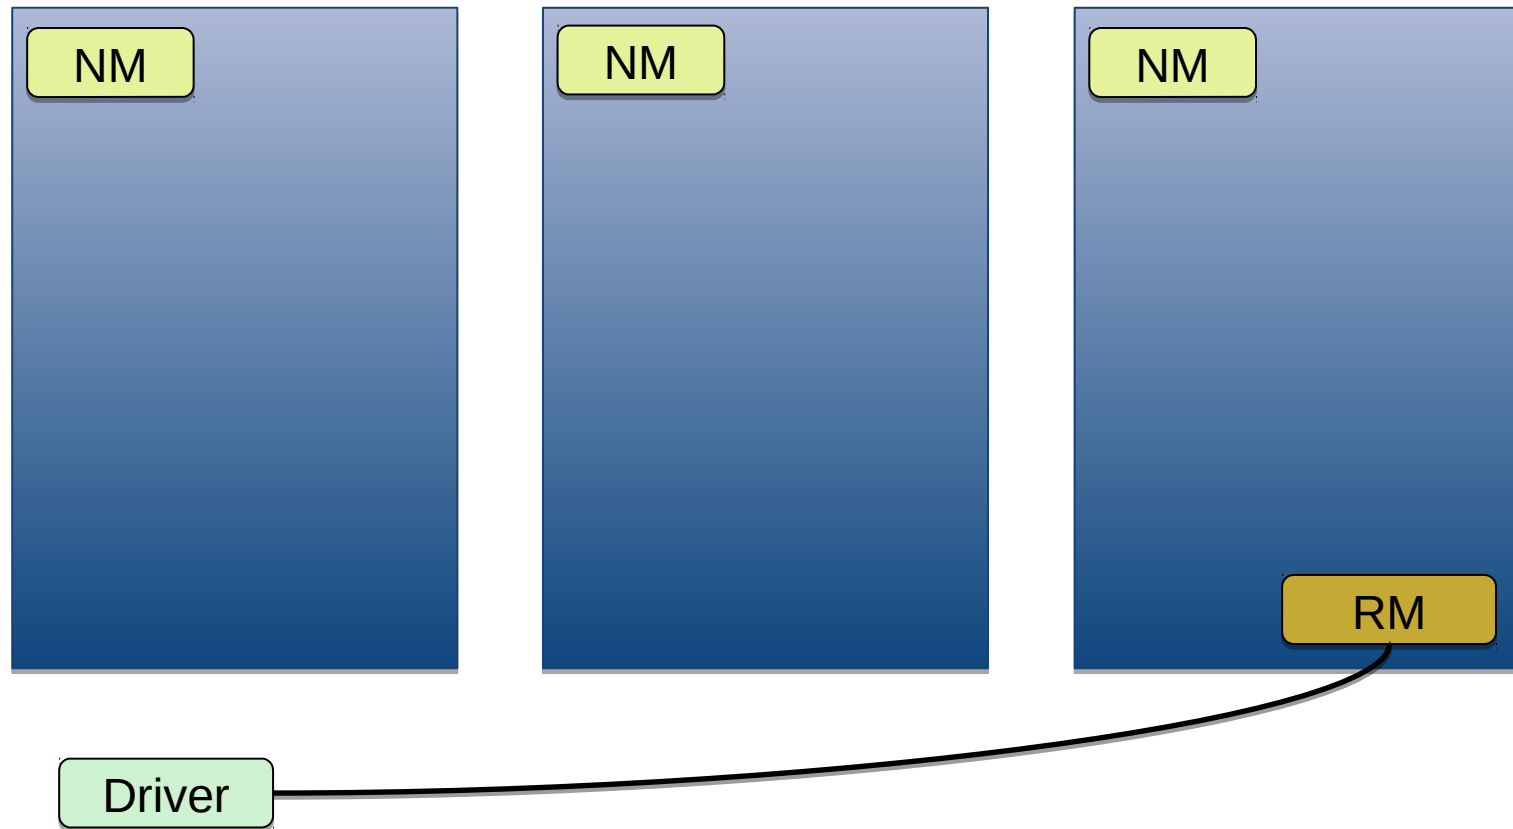
Spark Standalone



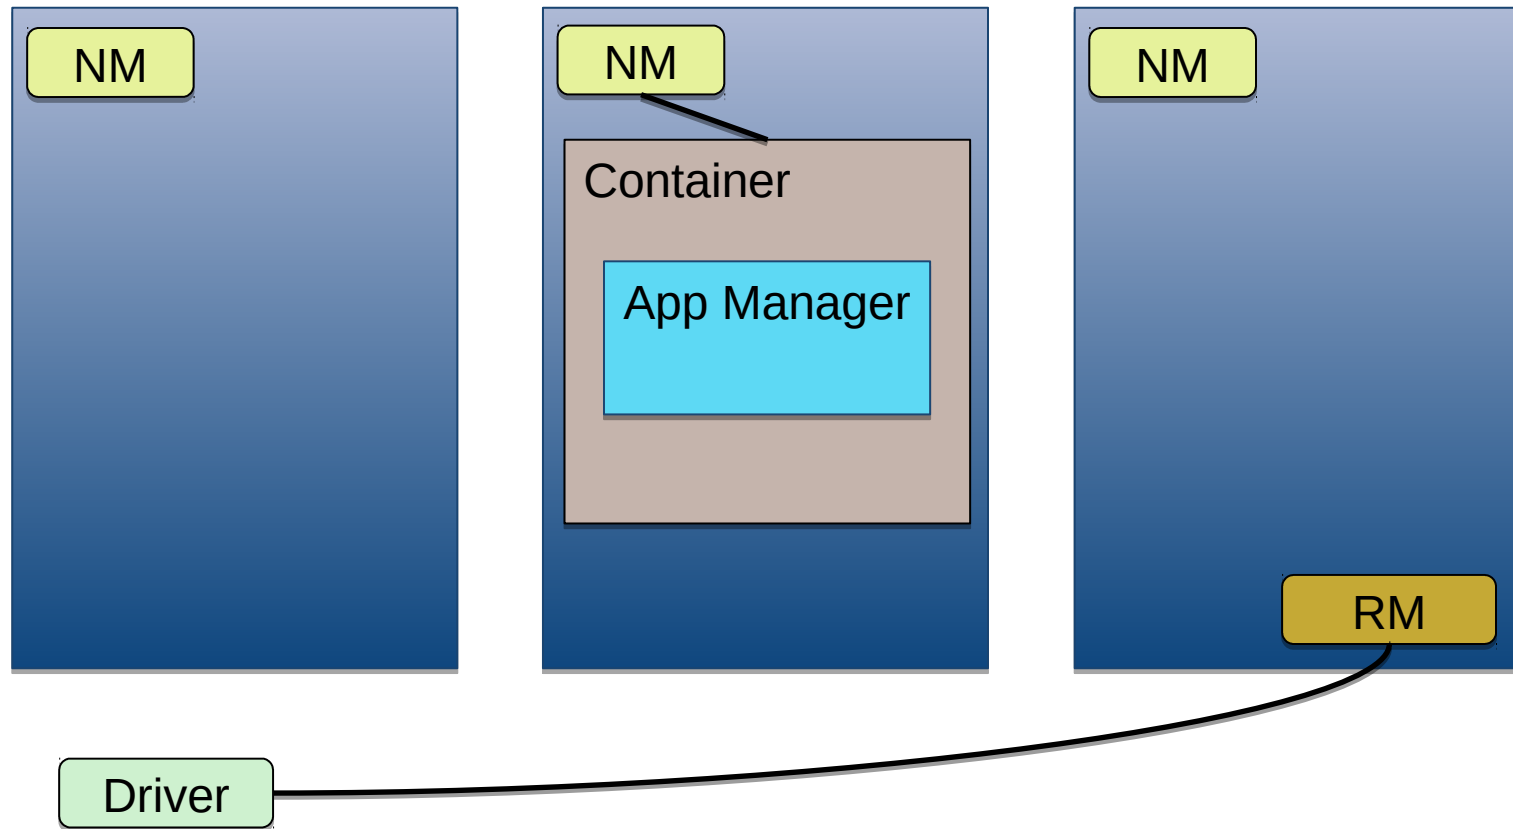
Spark YARN



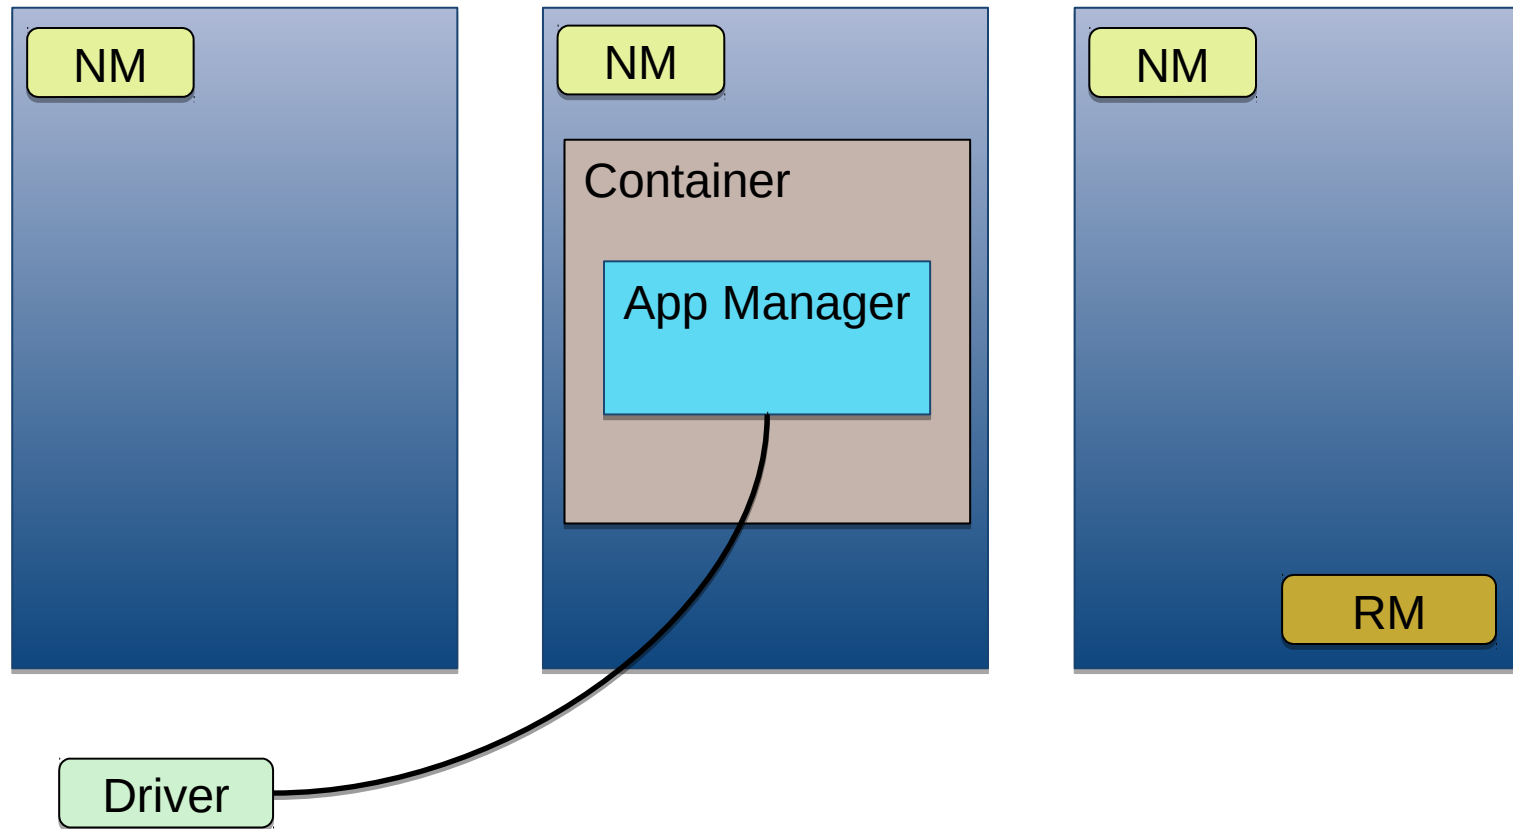
Spark YARN



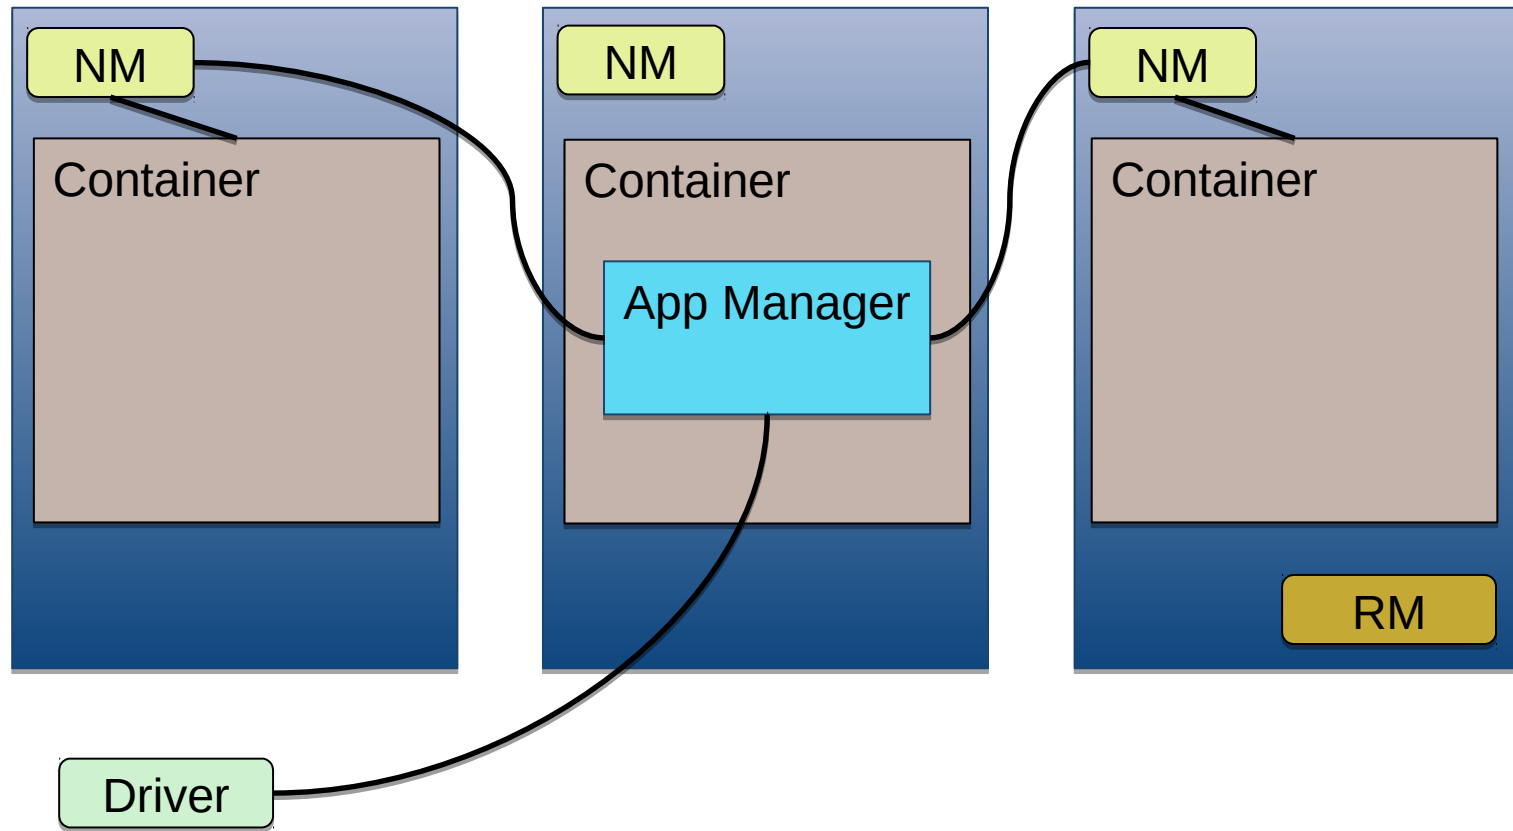
Spark YARN



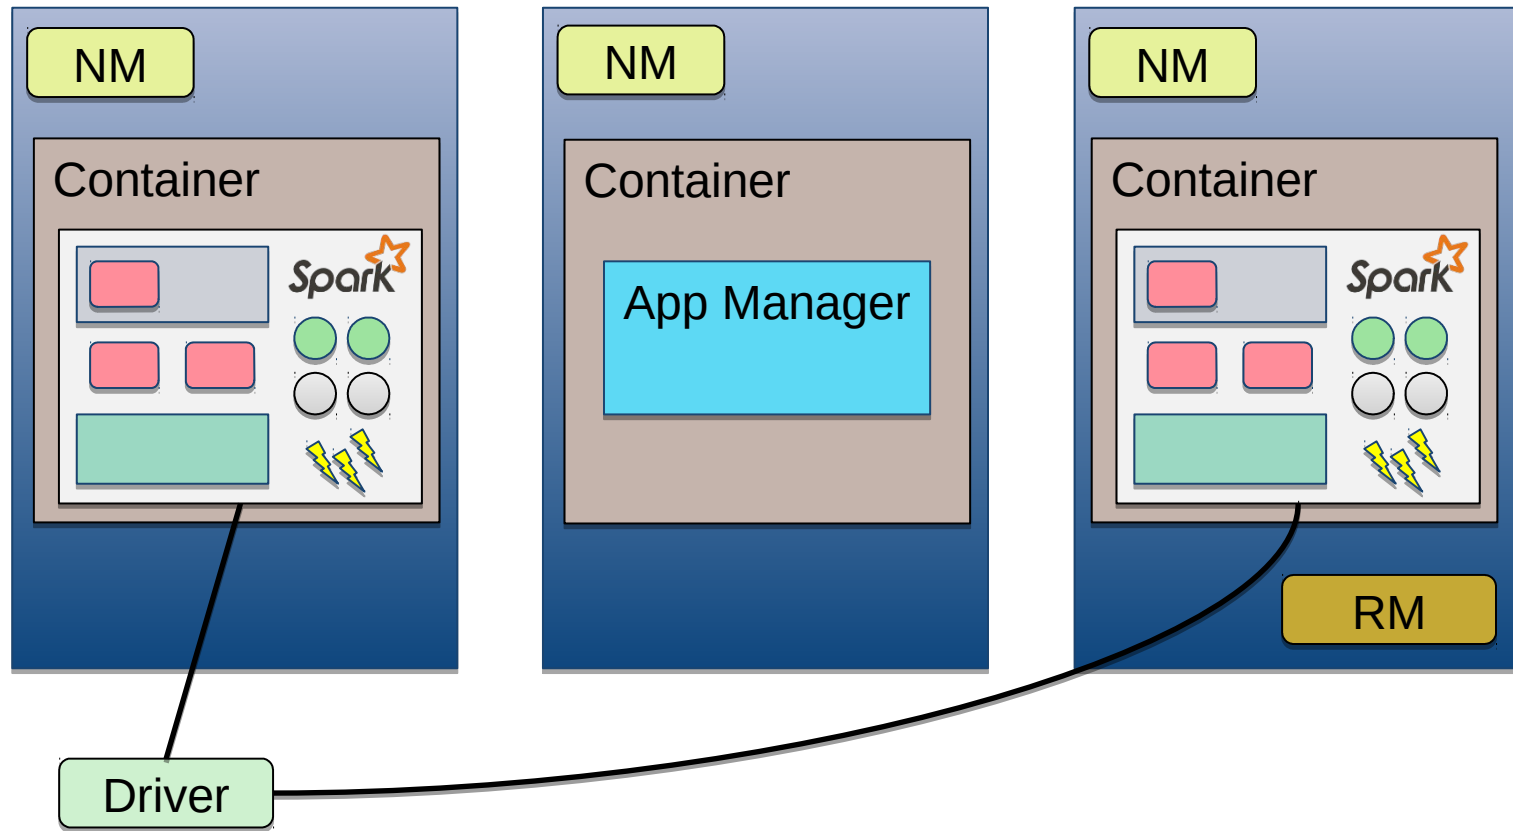
Spark YARN



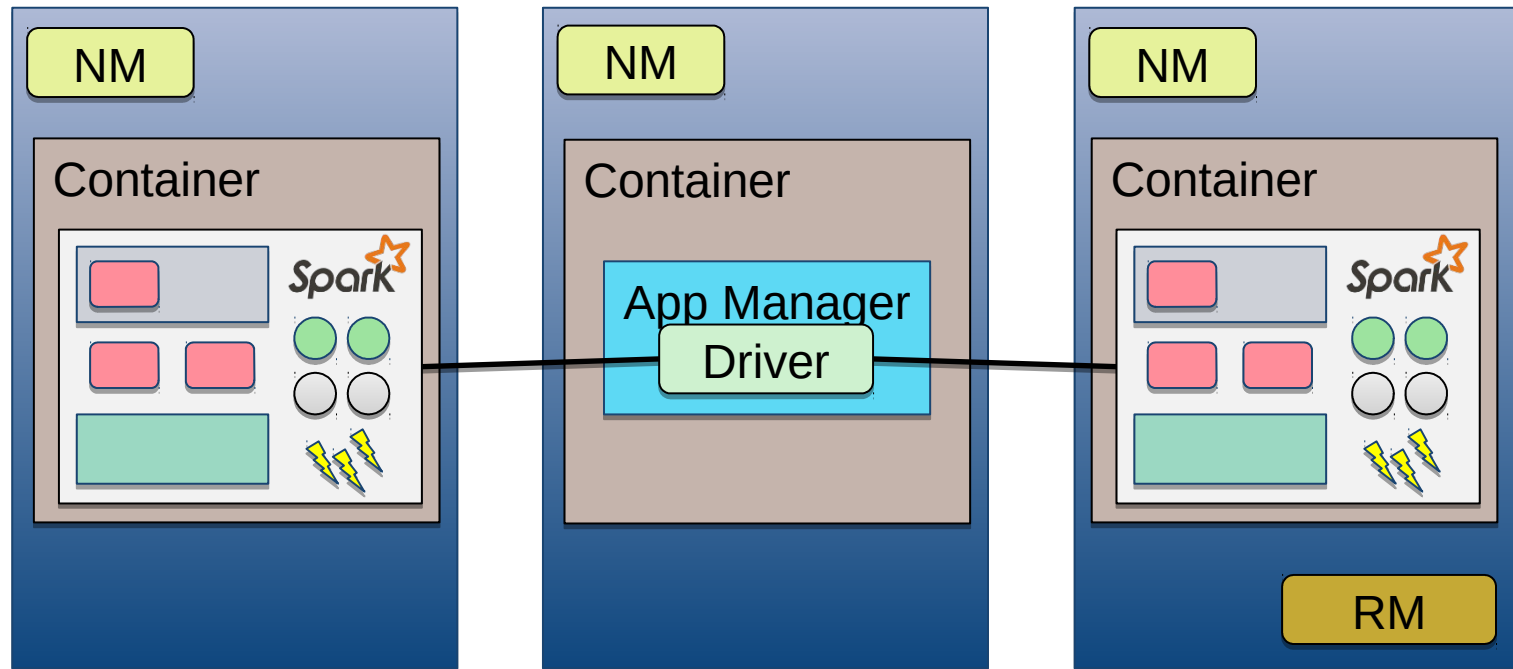
Spark YARN



Spark YARN (yarn-client)

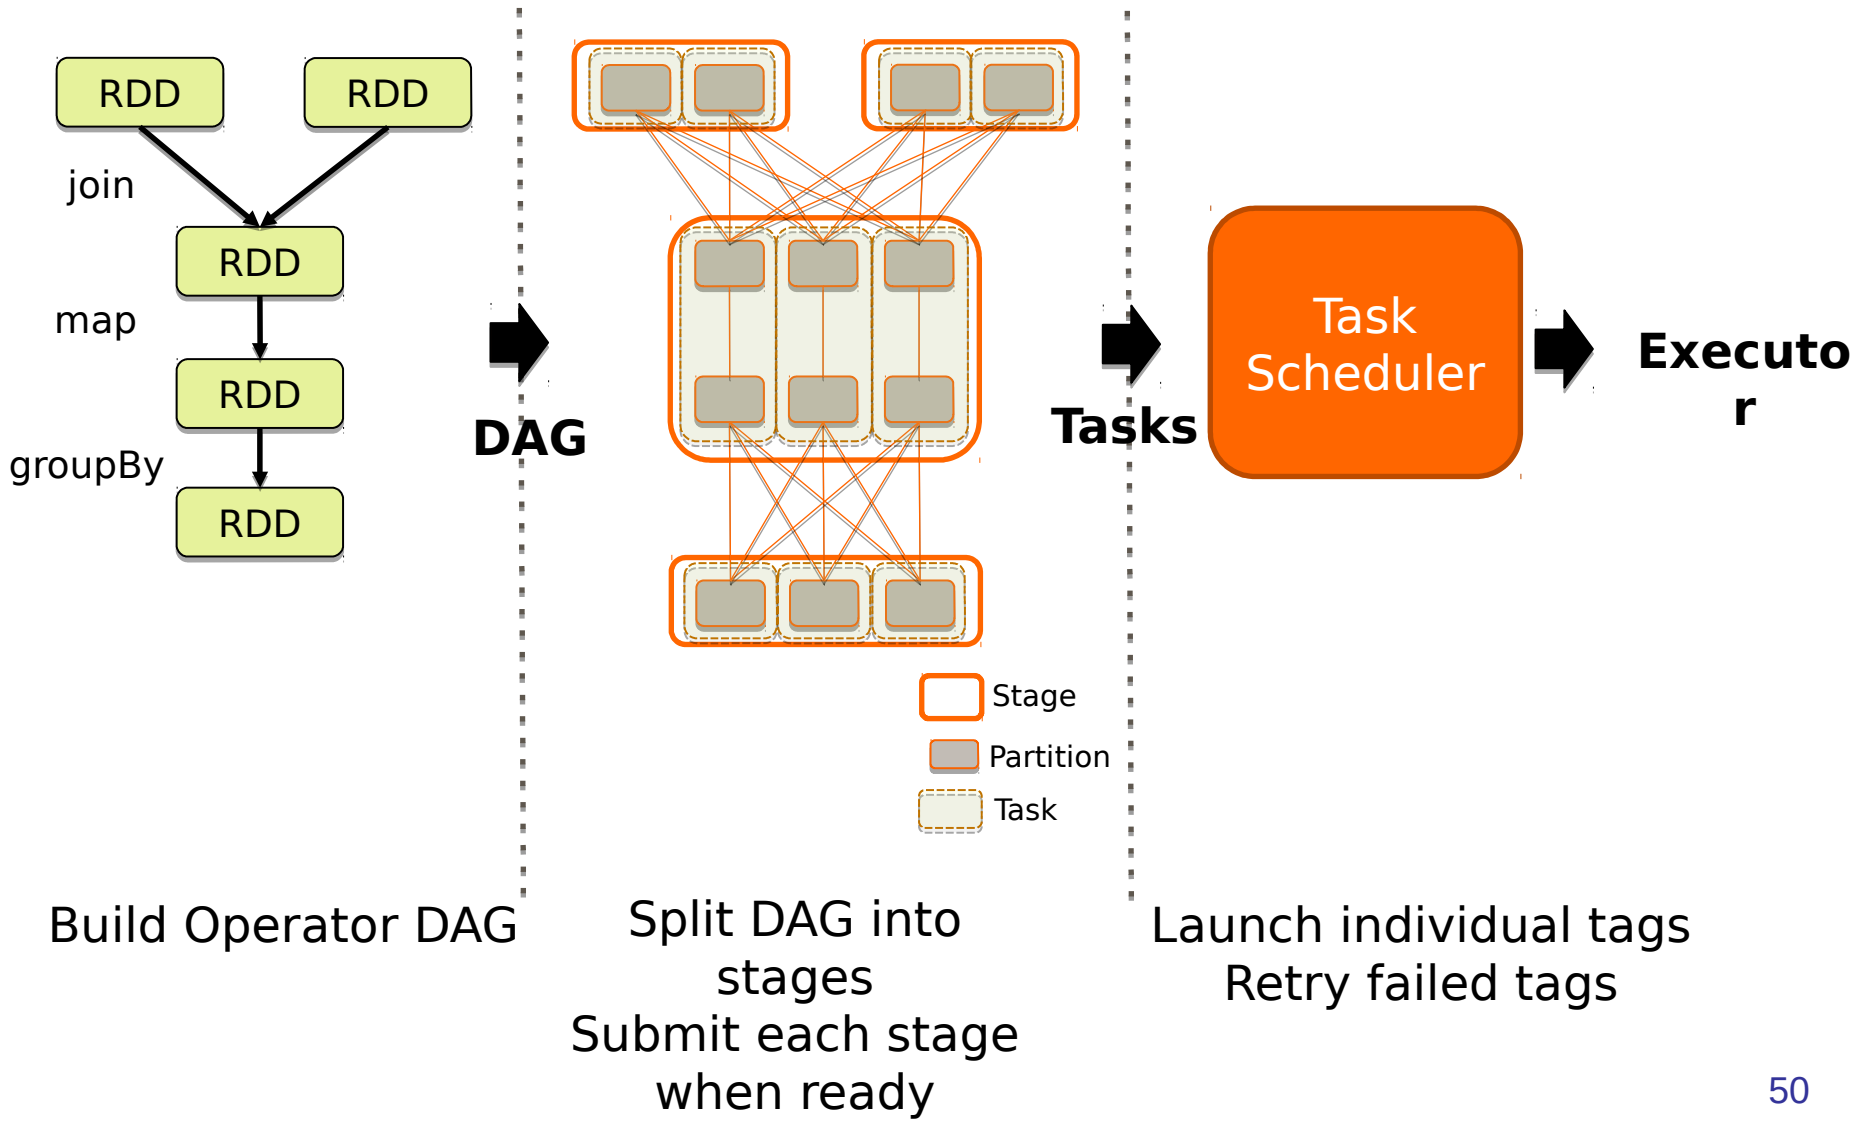


Spark YARN (yarn-cluster)



6. Operations

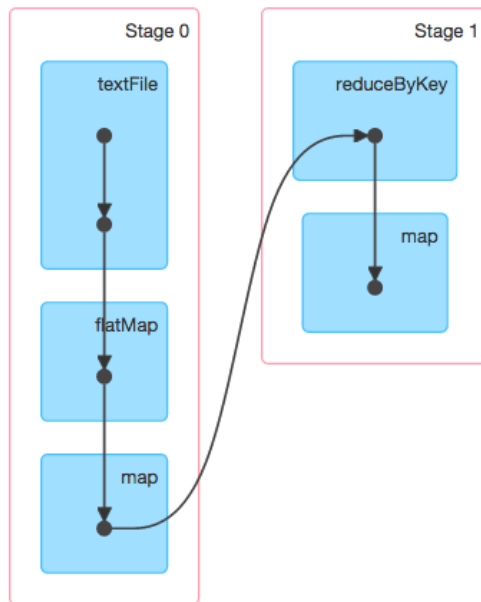
Scheduling



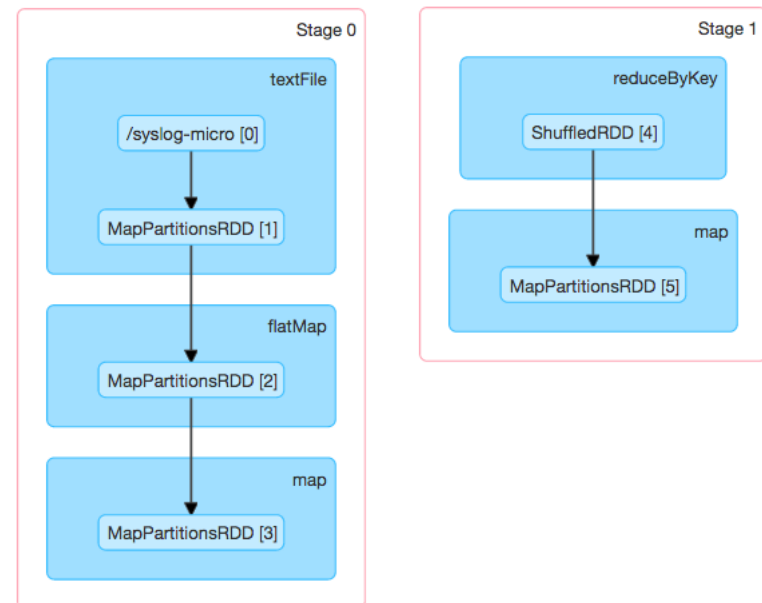
Scheduling

```
val logMessages = sc.textFile("hdfs://input/path")
val logEntries = logMessages.flatMap(Syslog.parseLine)
val wins = logEntries.map(logEntry => (wgs.winID(logEntry), new LogWinStats(logEntry)))
    .reduceByKey(_ + _)
    .map{ case (winID, winStats) => LogWindow(winID, winStats) }
wins.count()
```

DAG



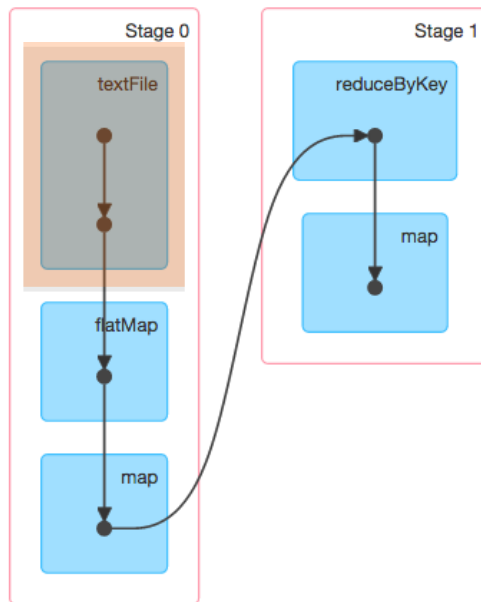
Stages



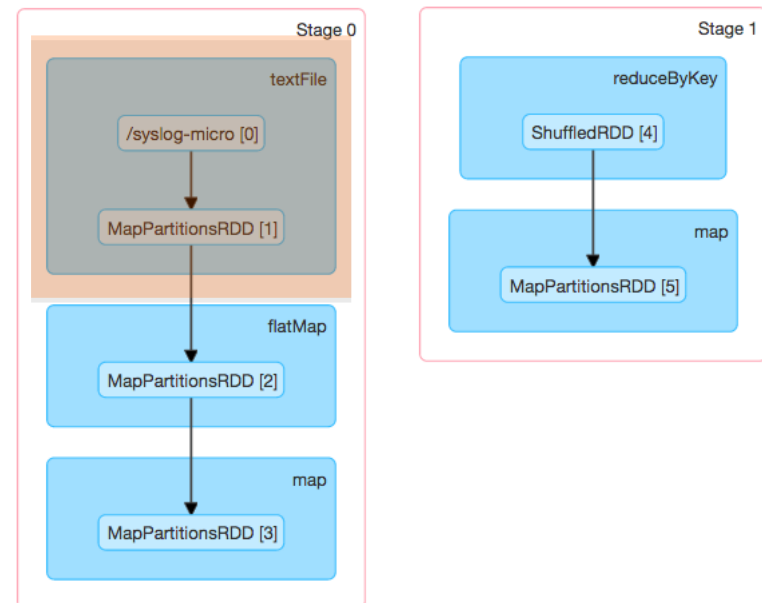
Scheduling

```
val logMessages = sc.textFile("hdfs://input/path")
val logEntries = logMessages.flatMap(Syslog.parseLine)
val wins = logEntries.map(logEntry => (wgs.winID(logEntry), new LogWinStats(logEntry)))
    .reduceByKey(_ + _)
    .map{ case (winID, winStats) => LogWindow(winID, winStats) }
wins.count()
```

DAG



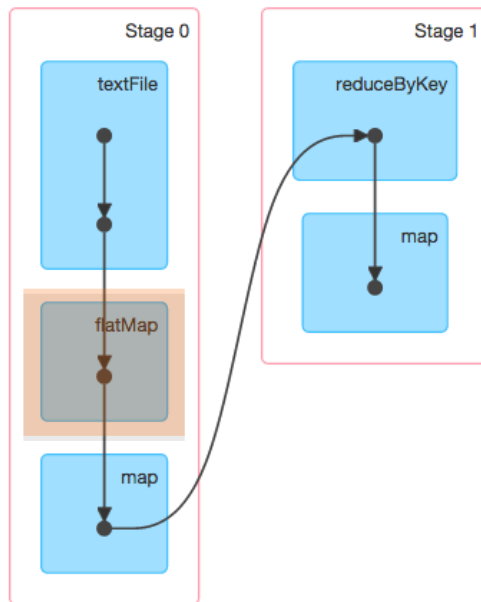
Stages



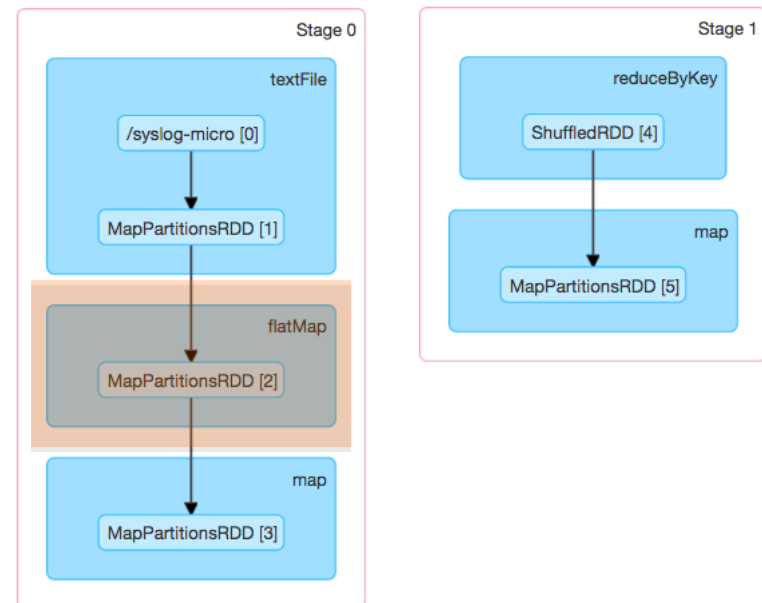
Scheduling

```
val logMessages = sc.textFile("hdfs://input/path")
val logEntries = logMessages.flatMap(Syslog.parseLine)
val wins = logEntries.map(logEntry => (wgs.winID(logEntry), new LogWinStats(logEntry)))
    .reduceByKey(_ + _)
    .map{ case (winID, winStats) => LogWindow(winID, winStats) }
wins.count()
```

DAG



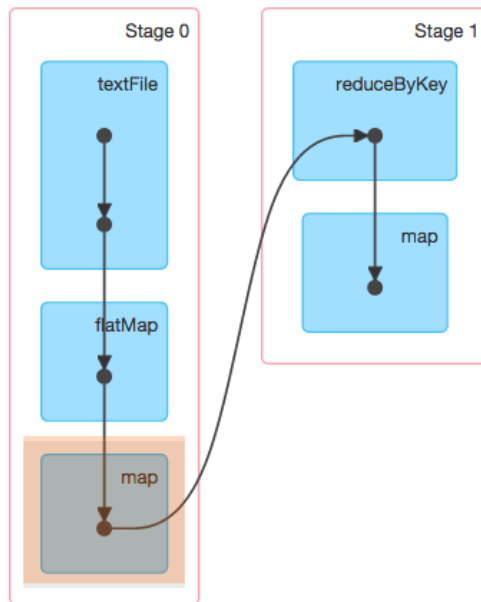
Stages



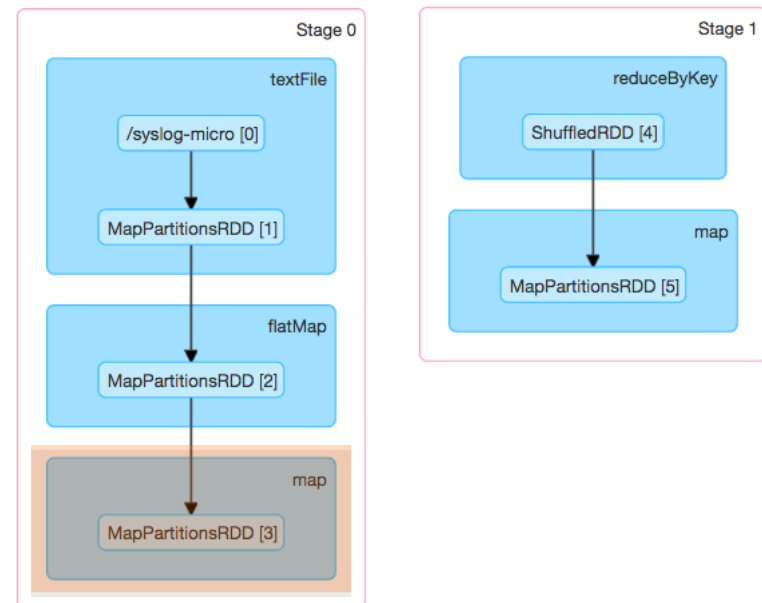
Scheduling

```
val logMessages = sc.textFile("hdfs://input/path")
val logEntries = logMessages.flatMap(Syslog.parseLine)
val wins = logEntries.map(logEntry => (wgs.winID(logEntry), new
LogWinStats(logEntry)))
    .reduceByKey(_ + _)
    .map{ case (winID, winStats) => LogWindow(winID, winStats) }
wins.count()
```

DAG



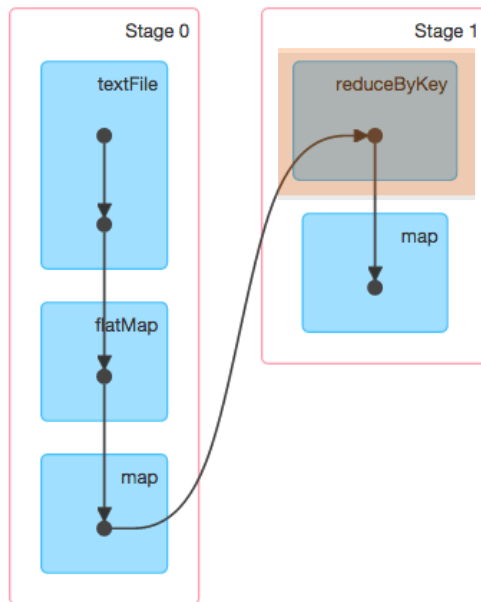
Stages



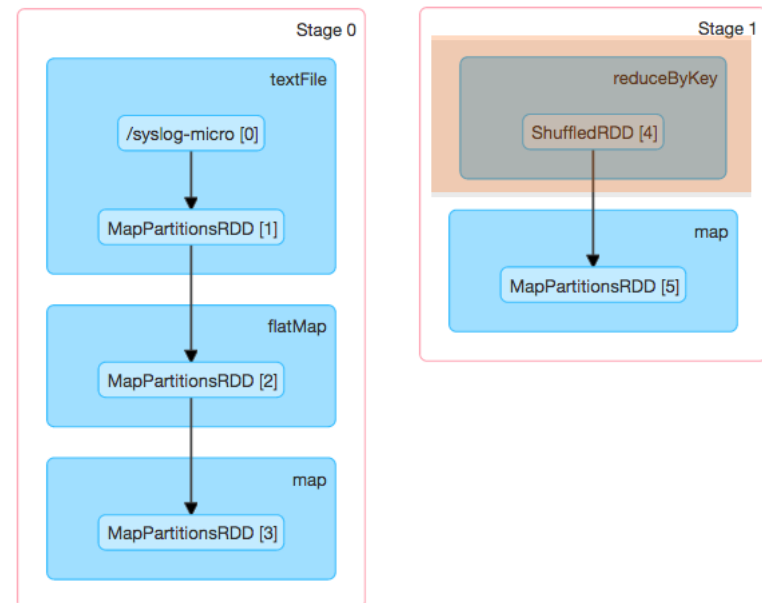
Scheduling

```
val logMessages = sc.textFile("hdfs://input/path")
val logEntries = logMessages.flatMap(Syslog.parseLine)
val wins = logEntries.map(logEntry => (wgs.winID(logEntry), new LogWinStats(logEntry)))
    .reduceByKey(_ + _)
    .map{ case (winID, winStats) => LogWindow(winID, winStats) }
wins.count()
```

DAG



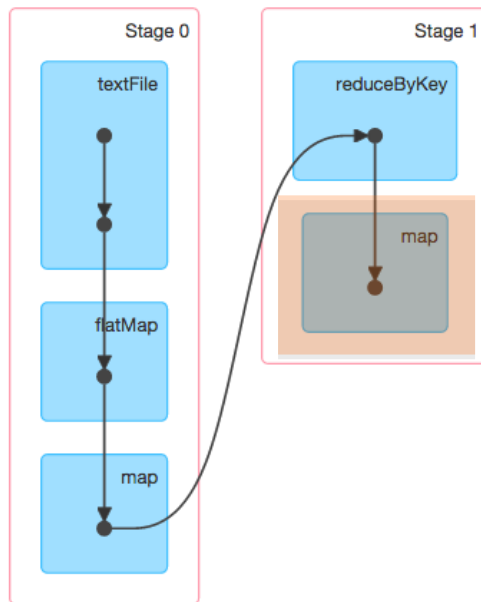
Stages



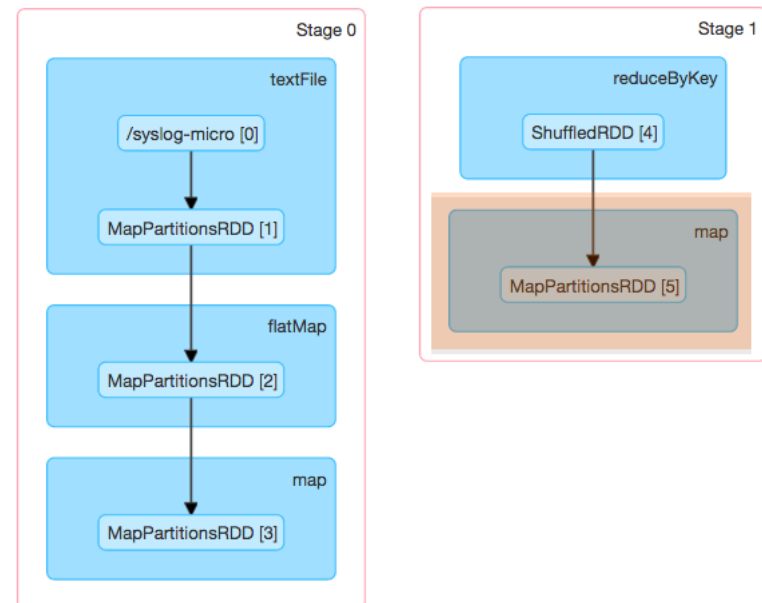
Scheduling

```
val logMessages = sc.textFile("hdfs://input/path")
val logEntries = logMessages.flatMap(Syslog.parseLine)
val wins = logEntries.map(logEntry => (wgs.winID(logEntry), new LogWinStats(logEntry)))
    .reduceByKey(_ + _)
    .map{ case (winID, winStats) => LogWindow(winID, winStats) }
wins.count()
```

DAG



Stages



Shuffle

- Represents a physical repartitioning of the data, for instance
 - Sorting a DataFrame
 - Joins
 - Aggregations
 - Grouping data that was loaded from a file by key
- Separates stages in the DAG
- May be triggered by
 - Repartitions
 - *ByKey and *By operations
 - Join
- It takes into account RDD partitioning
 - Narrow dependencies
 - Wide dependencies

Shuffle

- Sorting is only performed if required in reduce task
- *Hash Shuffle*
 - Each map creates one file per reducer
 - Records are appended to each file
- *Sort shuffle (default from 1.4)*
 - One file per map task
 - File ordered by reducer id and indexed
- *Tungsten shuffle*
 - Operates directly on serialize data
 - To be used only under certain conditions

Hash Shuffle

- Each map creates one file for each reducer
 - Files = $M * R$
- Each record is written directly to target file
- Optimization – File consolidation
 - A pool of output files created by executor
 - Each map tasks requests R files
- Pros / Cons
 - ++ fast / no sort, no memory overhead, no hash table
 - -- # partitions \propto a lot of files, random IO

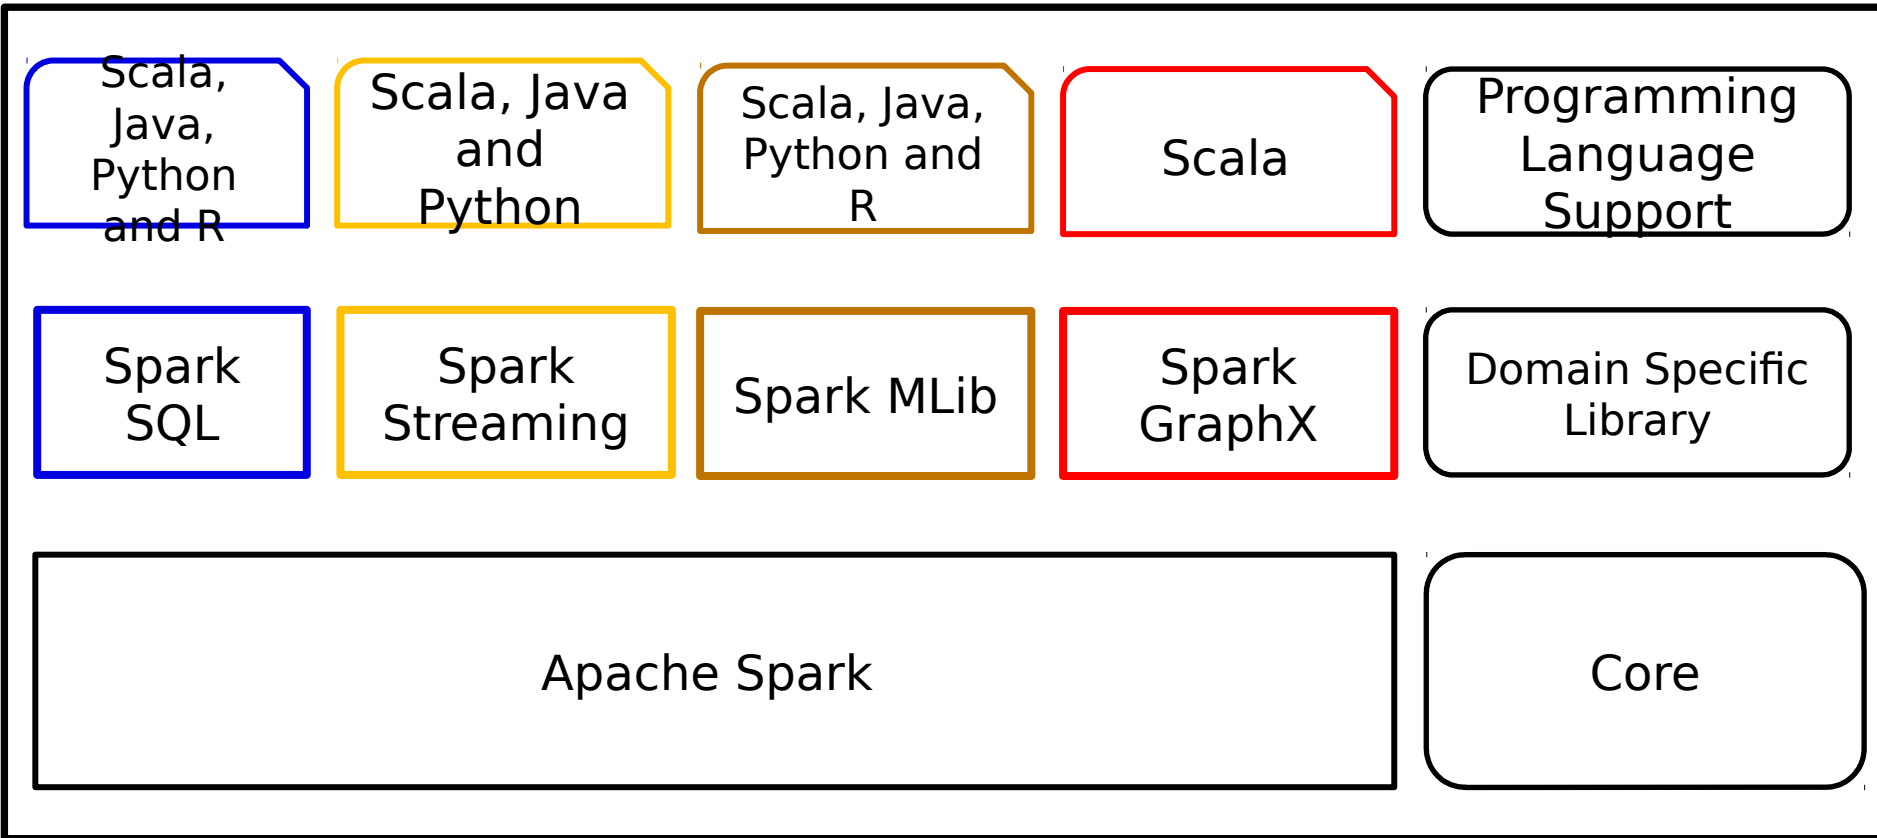
Sort Shuffle

- Similar to MapReduce shuffle
 - One file per map task
 - File ordered by reducer id and indexed
 - If not enough memory \Rightarrow spill to disk
 - BUT results are not merged in reduce side
- Special implementation of hash table
 - Allows combine in place
- Merging is only done when request from reducerPros / Cons
 - ++ less # of files, less # of random IO
 - -- sorting slower than hashing, SSDs favor hash-based

Tungsten Shuffle

- Operates directly on serialized binary data
 - Direct spilling
 - May merge spills by direct concatenation
- Uses special cache-efficient sorter
- To be used only when
 - No aggregation
 - Codec supports relocation of serialized values
 - Less than 16777216 partitions
 - Record < 128MB

7. Tools



SparkSQL

- Provides the ability to run SQL on top of Spark
- A Spark module for structured data processing
- Defines an interface for a semi-structured data type called Dataframes and a typed version (Datasets)
 - APIs as well support for basic SQL queries
- The Spark SQL library enables
 - Relational processing both within Spark programs and on external data sources using a programmer-friendly API
 - High performance using established DBMS techniques
 - Support of new data sources, including semi structured data and external databases
 - Enables extension with advanced analytics algorithms such as graph processing and machine learning

SparkSQL Example

Create a DataFrame, manipulate it with SQL, and then manipulate it again as a DataFrame

in Python

```
# spark is an existing SparkSession
    spark.sql("CREATE TABLE IF NOT EXISTS src (key INT, value
STRING) USING hive")
    spark.sql("LOAD DATA LOCAL INPATH
'examples/src/main/resources/kv1.txt' INTO TABLE src")

# Queries are expressed in HiveQL
    spark.sql("SELECT * FROM src").show()
# +---+-----+
# |key|  value|
# +---+-----+
# |238|val_238|
# | 86|val_86|
# |311|val_311|
```


Dataframes

- Structured API which represents a table of data with rows and columns
- Distributed collection of data organized into named columns
- DF Creation
 - From existing RDD
 - Schema inference
 - Schema specification
 - Data sources
 - Parquet files, JSON, (csv)
 - Hive, JDBC
- DF Manipulation
 - DF Operations
 - SQL queries programmatically

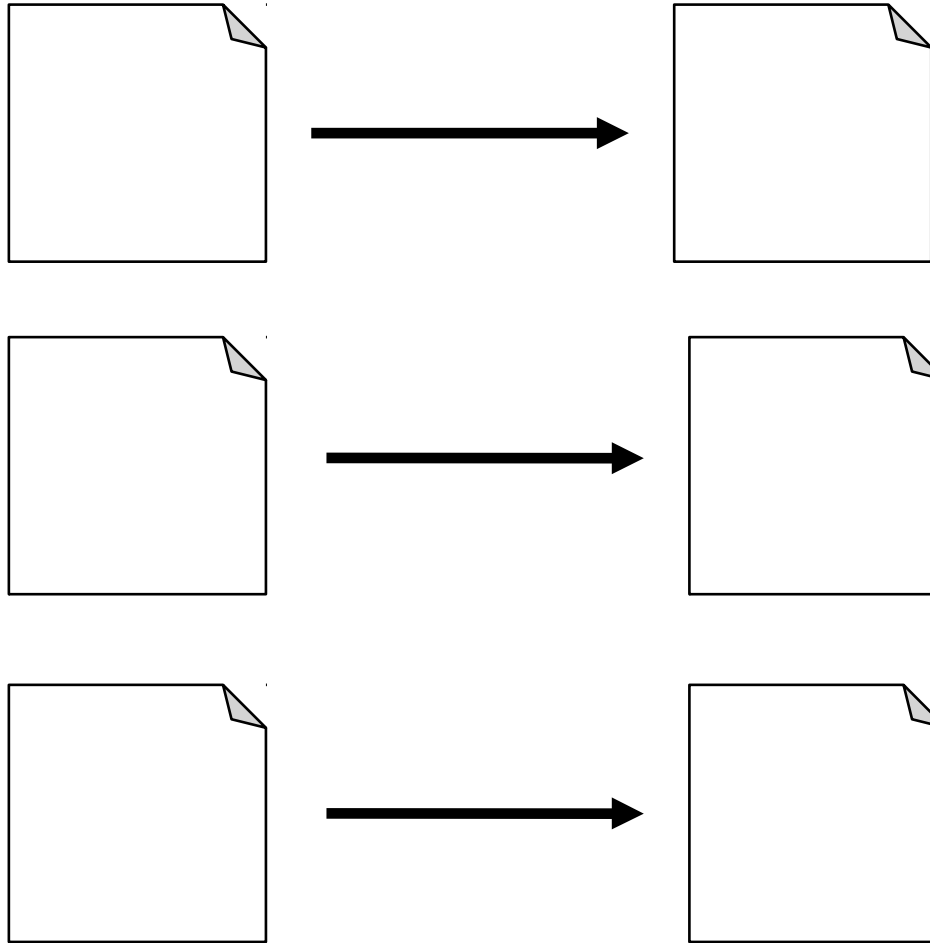
Dataframes

- Dataframe represents
 - Immutable, lazily evaluated plans
- Specify what operations to apply to data residing at a location to generate some output
- Immutable
 - Spark core data structure cannot be changed after they are created
- Lazy evaluation
 - Spark will wait until the last moment to execute the graph of computation instructions

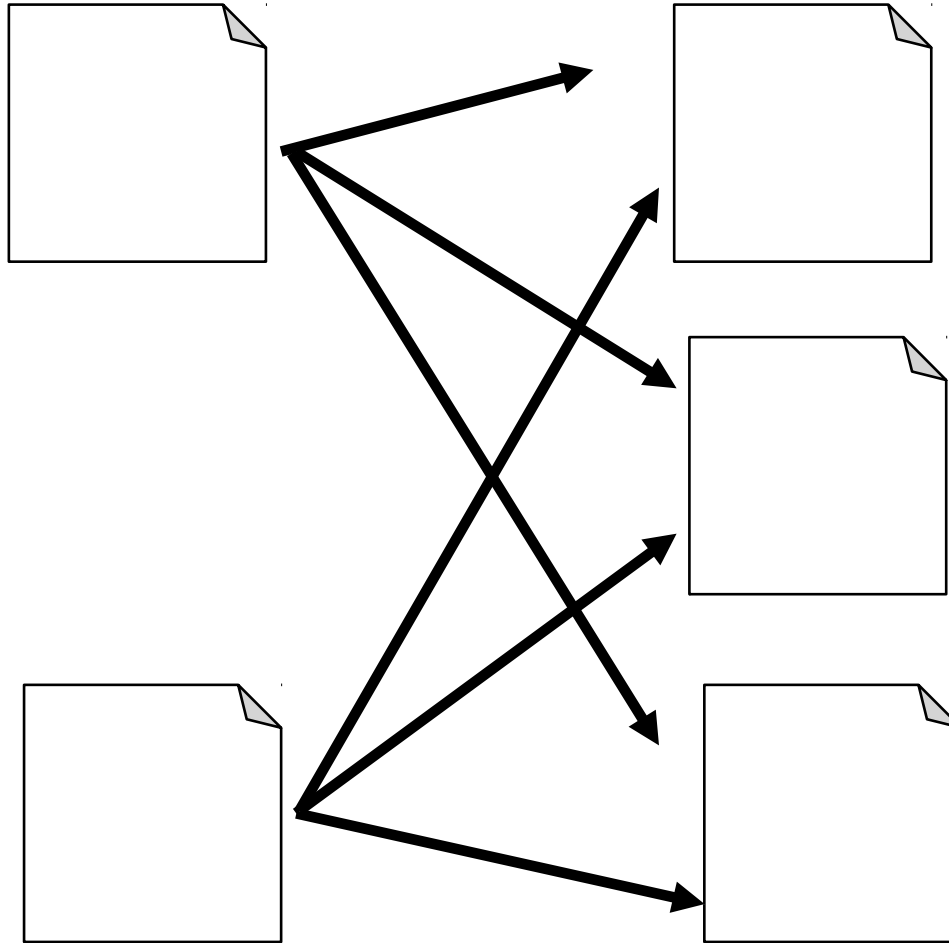
Dataframe operations

- Transformations: ways of specifying different series of data manipulation.
 - Build up logic transformation plan
 - Narrow dependencies (each input partition will contribute to only one output partition)
 - Wide dependencies (input partitions contributing to many output partitions)
- Action : instructs Spark to compute a result from a series of transformations (e.g. Count)
 - Actions to view data in the console
 - Actions to collect data to native objects in the respective language
 - Actions to write to output data sources

Narrow transformations 1 to 1



Wide transformations 1 to N



Dataframe Example

Schema inference

```
case class LogEntry(  
  timestamp: Long,  
  host: String,  
  facility: Int,  
  severity: Int,  
  app: String,  
  message: String)
```

```
val logMessages = sc.textFile("hdfs://input/path")  
val logEntries = logMessages.flatMap(Syslog.parseLine)  
val logEntriesDF = logEntries.toDF()
```

Dataframe Example

Schema inference

```
logEntriesDF.printSchema()
```

```
root
```

```
|-- timestamp: long (nullable = false)
```

```
|-- host: string (nullable = true)
```

```
|-- facility: integer (nullable = false)
```

```
|-- severity: integer (nullable = false)
```

```
|-- app: string (nullable = true)
```

```
|-- message: string (nullable = true)
```

Dataframe Example

Querying with DF API

```
val res = logEntriesDF.filter("severity < 3")  
                        .groupBy("app")  
                        .count()  
                        .filter("count > 5")  
                        .sort($"count".desc)
```

```
res.show(5)
```

```
+-----+-----+  
|      app|count|  
+-----+-----+  
|   kernel| 1942|  
|lustre-target| 514|  
|  OMPI-ERROR| 331|  
|    ioadm| 297|  
|    crmd|  18|  
+-----+-----+
```


Dataframe Example

Querying with SQL

```
logEntriesDF.registerTempTable("logEntries")
val res = sqlContext.sql(
  "select app, count(*) as count from logEntries where severity < 3 " +
  "group by app having count > 5 " +
  "order by count desc"
)
res.show(5)
```

app	count
kernel	1942
lustre-target	514
OMPI-ERROR	331
ioadm	297
crmd	18

Dataframe Example

Programmatically specifying a schema

```
val schema = StructType(Seq(  
  StructField("logCount", LongType),  
  StructField("facilityCount", ArrayType(LongType)),  
  StructField("severityCount", ArrayType(LongType)),  
  StructField("appCount", MapType(StringType, LongType))  
))  
  
val windowsRow = windows.map(w => Row(w.stats.logCount,  
                                       w.stats.facilityCount,  
                                       w.stats.severityCount,  
                                       w.stats.appCount))  
  
val windowsDF = sqlContext.createDataFrame(windowsRow, schema)  
windowsDF.printSchema()
```

Dataframe Example

Programmatically specifying a schema

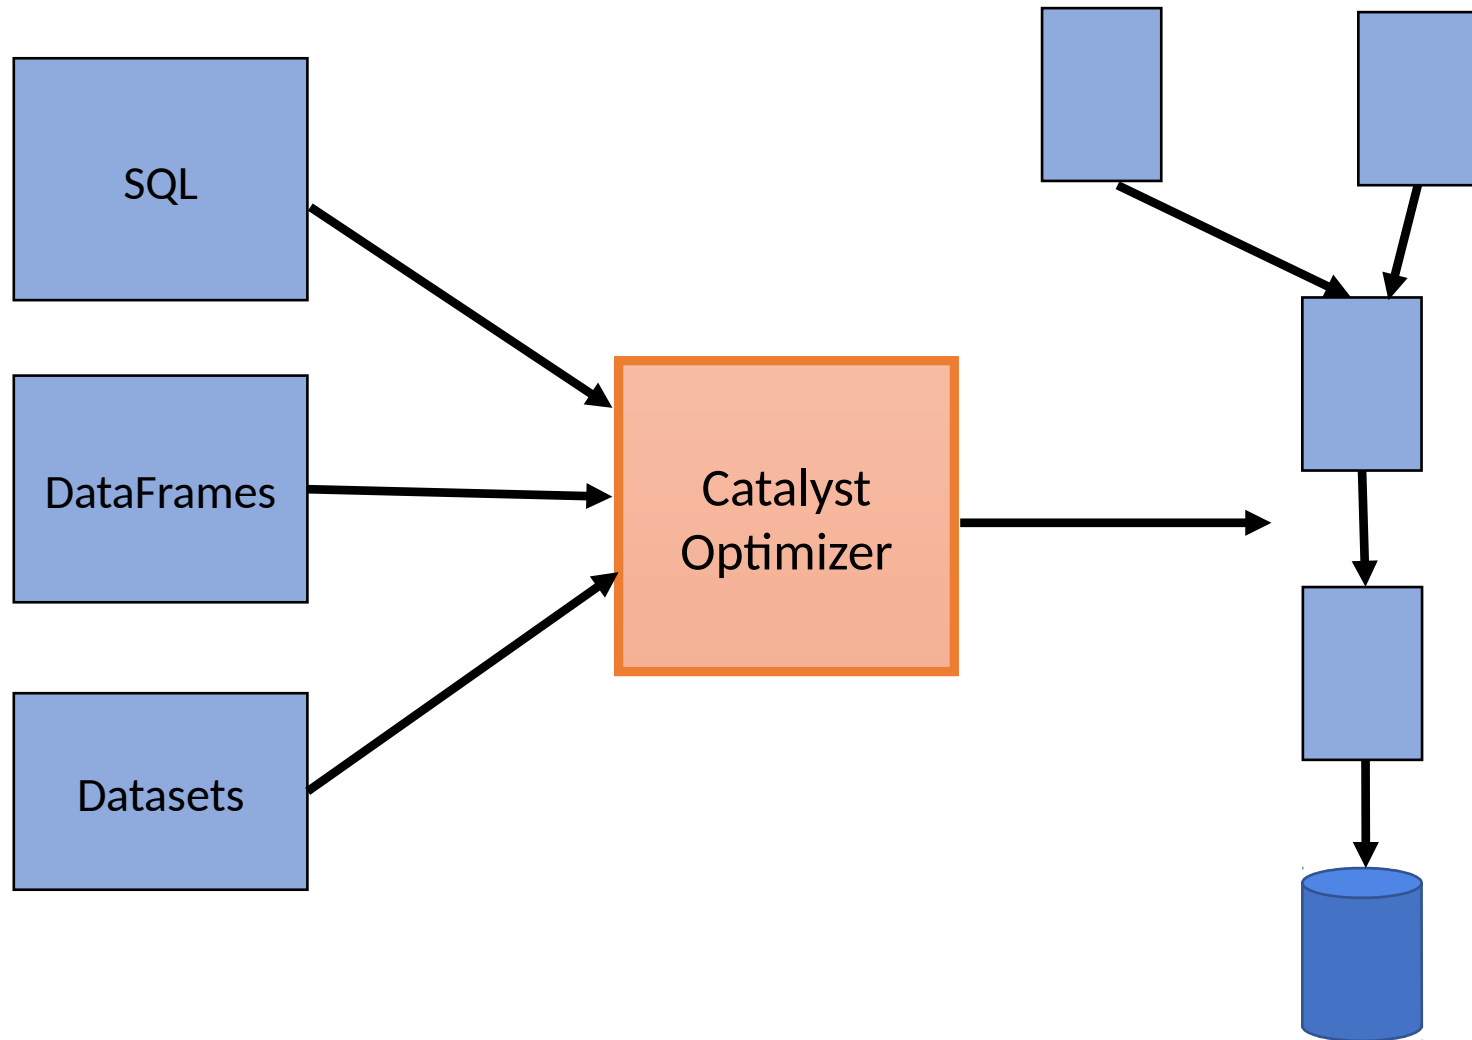
windowsDF.printSchema()

root

```
|-- logCount: long (nullable = true)
|-- facilityCount: array (nullable = true)
|   |-- element: long (containsNull = true)
|-- severityCount: array (nullable = true)
|   |-- element: long (containsNull = true)
|-- appCount: map (nullable = true)
|   |-- key: string
|   |-- value: long (valueContainsNull = true)
```

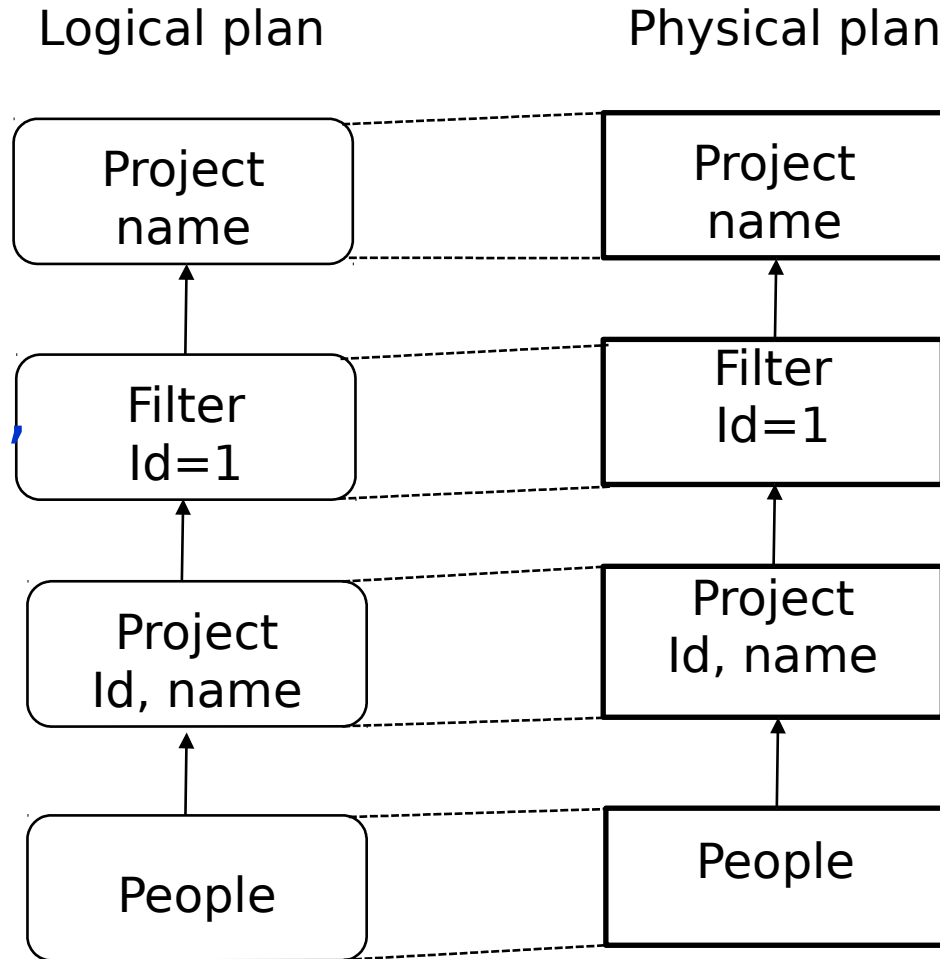
Dataframe Optimization

Catalyst optimizer



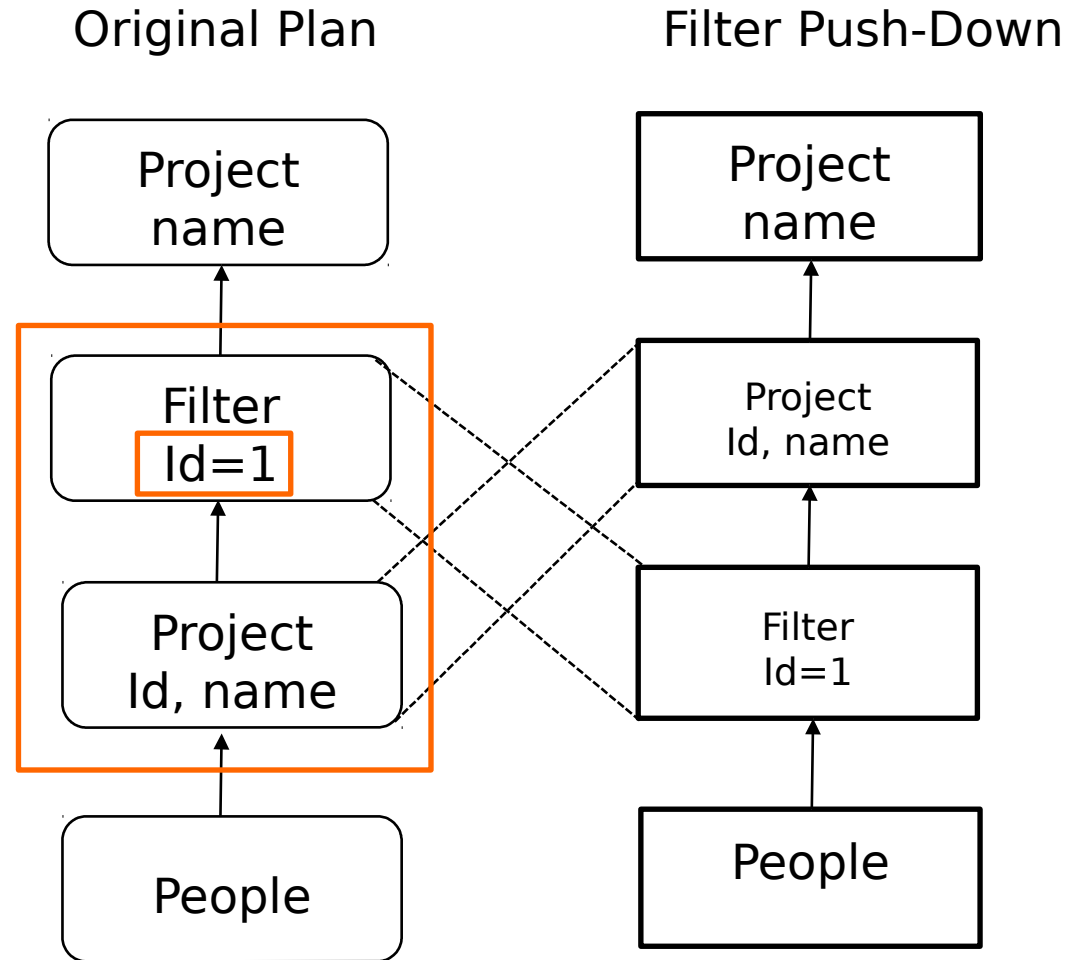
Spark SQL Optimization

SELECT name
FROM (SELECT id,
name
FROM People) p
WHERE p.id = 1



Spark SQL Optimization

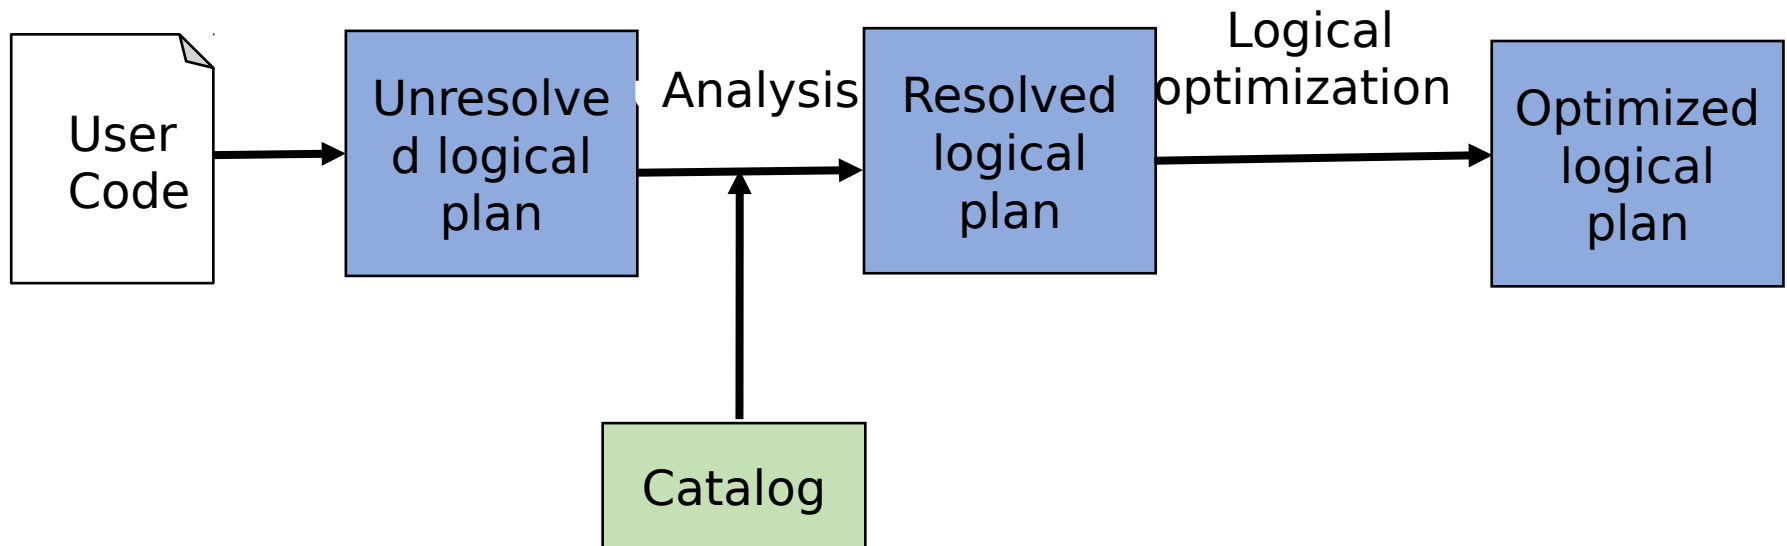
1. Find filters on top projections
2. Check that the filter can be evaluated without the result of the project
3. If so, switch the operators



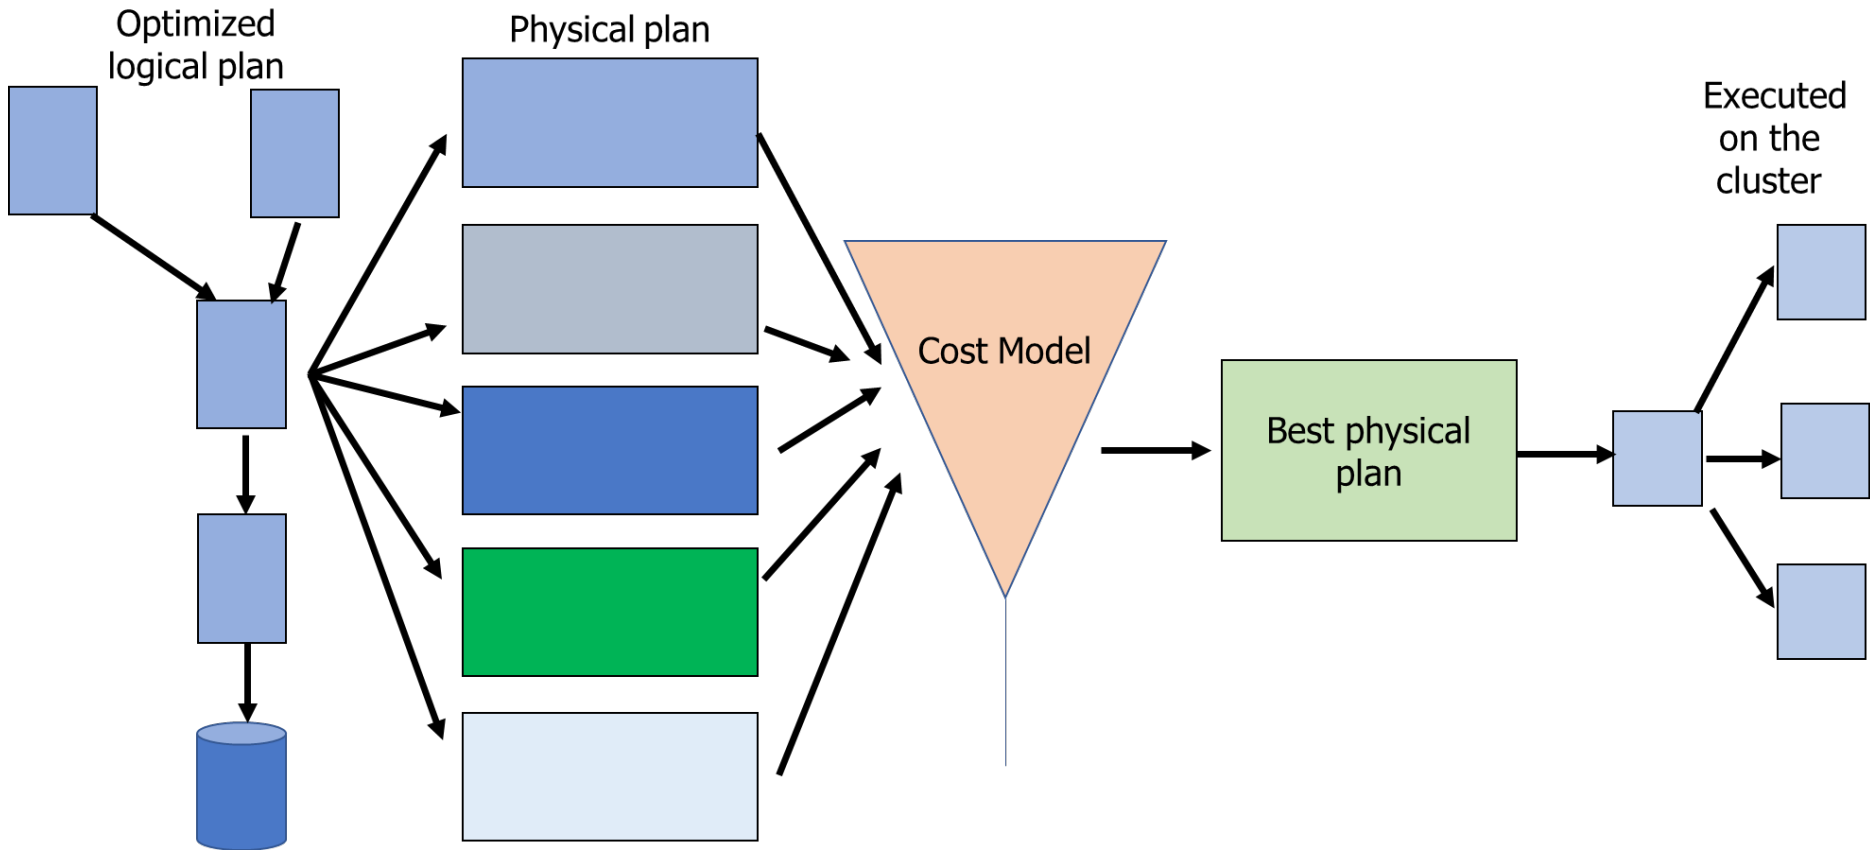
Dataframe Optimization

- Catalyst optimizer [SIGMOD 2015]
 - Supports rule-based and cost-based optimizations
 - Users can extend optimizer
 - Data source specific rules
 - e.g., push filtering/aggregation into external storage
- Spark performs standard optimization
 - Constant folding
 - Predicate pushdown
 - Projection pruning
 - Null propagation
 - Boolean expression simplification
 - etc.

Dataframe Optimization



Physical Planning



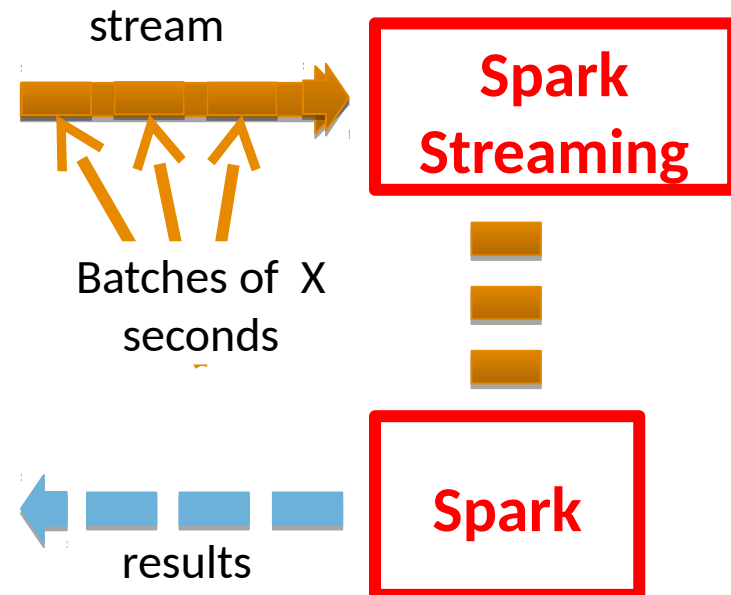
Spark Streaming

- Library that is working on top of Spark, that enables processing of live data streams
- Stream processing API
 - Data divided into mini-batches
- DStream \Rightarrow Sequence of RDDs
- Sources: HDFS, Kafka, Flume, Kinesis
- Supported languages
 - Scala
 - Java
 - And Python

Micro Batch Data Processing

Decomposition of a stream in a series of very small jobs (*mini-batches*)

- Decomposition in batches of X seconds (ex. $X=0,5$)
- Each batch is treated as an RDD
- The results of RDD operations are returned as batches



SparkMLlib/ ML

- Enables development of machine learning applications
 - Classification and regression (linear models, decision trees)
 - Collaborative filtering (ALS)
 - Clustering (k-means, Gaussian mixture, LDA)
 - Dimensionality reduction (SVD, PCA)
 - etc
- Two packages
 - MLlib: on top of RDDs
 - ML: on top of DataFrames
- Supporting languages
 - Scala
 - Java
 - Python
 - And R

ML API

- Based on DataFrames
- ML workflows can be specified as pipelines
 - Pipelines can be trained with combination of parameters
- Pipeline elements
 - Transformer
 - Estimator
 - Evaluator

8. Conclusion on Spark

- Rich APIs to make data analytics *fast*
- Up to 100x speedups in real applications
- Framework for large-scale parallel processing
- Multiple integrated components

Sources

- Documents

- Apache Spark. <http://spark.apache.org/>
- Big Data Processing with Apache Spark. Miguel Liroz Gistau
- Parallel Programming with Spark. Matei Zaharia
- Spark/Python Practice. Boyan Kolev and Esther Pacitti

- Books

- Principles of Distributed Database Systems 3rd Edition, Springer. Tamer Özsu and Patrick Valduriez
- Spark The Definitive Guide: big data processing made simple 1st Edition, O'Reilly Media. Bill Chambers and Matei Zaharia
- High Performance Spark 1st Edition, High Performance Spark 1st Edition. Holden Karau and Rachel Warren
- Apache Spark 2 for Beginners, Coderprog. Rajanarayanan Thottuvaikkatumana



Spark Practice

1. Download and run Spark's Python shell
2. Basic exercise
3. Advanced exercise

1. Download and run Spark's Python shell

- Download Spark binaries from <https://spark.apache.org/downloads.html>
 - Choose some « pre-built for Apache Hadoop » package type
 - Or directly from the Apache mirror
- Extract the archive, enter into the extracted directory and start Spark's Python shell by running the `./bin/pyspark` executable
- Note that the number of workers can be passed with the `--master` argument of `pyspark`
 - E.g. This command initiates a Spark cluster with 2 workers : `./bin/pyspark --master local [2]`

2. Basic Exercise

- Word count
- Evaluating the benefit of parallelism
- Filtering the result

3. Advanced Exercise

- First goal: get user-keyword pairs
- Count the occurrence of each pair
- Map by keywords
- Final goal: keep only the winner
- Raising the scale
- Filter by particular keywords